

Defending Against Standard-Threat-Model Explanation Poisoning via Causal Proof of Reasoning

Ayush Chadha, Ragav Palaniswamy, Sanjana Gummuluru

Guide: Dr. M. Sridevi

National Institute of Technology Tiruchirappalli

Abstract—Federated Learning (FL) has emerged as a foundational paradigm for privacy-preserving machine learning. Recent advances extend FL toward *agentic* systems, where autonomous clients independently curate data, optimize objectives, and adapt learning strategies. In parallel, Explainable AI (XAI) has been adopted as the primary mechanism for trust and regulatory compliance. However, recent work demonstrates that XAI itself can be adversarially manipulated, enabling *Explanation Poisoning*: attacks that preserve accuracy and plausible explanations while embedding hidden backdoors.

This work proposes a novel governance framework termed *Proof of Reasoning (PoR)*. Instead of auditing numerical model updates, PoR enforces structural alignment of causal reasoning across agents using Graph Edit Distance (GED). We show that explanation-poisoning attacks are fundamentally incompatible with causal-graph consistency, creating a mathematical contradiction that forces adversaries to either abandon stealth or abandon the attack. We contribute: (1) a Dual-Gate protocol combining structural GED and Causal Effect Divergence (CED) checks with robust MAD calibration; (2) a five-adversary stress-testing pool (Temporal Mimicry, Topology Inversion, Gradient Mimicry, Adaptive RL Evasion, Weight-Only Coefficient Manipulation) plus a Sybil attack framework; (3) a Bayesian Dirichlet-Multinomial consensus mechanism with sliding-window credibility decay and cross-client correlation penalties; (4) a curriculum grace period with robust coordinate-wise median aggregation for liveness preservation; (5) graph differential privacy with correct single-mechanism Laplace calibration; (6) a production-grade engineering infrastructure including CI/CD pipeline, pre-commit quality assurance, Makefile-driven workflow, and pre-flight configuration validation. We provide theoretical justification (Theorems 1 and 1.1 with formal proofs), architectural design, a comparative analysis against three generations of defenses, and empirical validation across two structurally distinct domains: (i) the ASIA Bayesian network [40], where PoR detects adversaries in 9 of 10 federated rounds, and (ii) a highly non-stationary Algorithmic Finance domain using a Transformer Actor-Critic agent [24] with PPO [28], DP-SGD [20], and GARCH(1,1) market simulation [32], where PoR achieves 100% detection of Temporal Mimicry and Topology Inversion adversaries at a 2.8% honest false-positive rate over 15 federated rounds. The framework is validated across a 39-test suite covering all core components.

I. INTRODUCTION & MOTIVATION

A. The Evolution: From Passive Learning to Agentic Autonomy

The field of distributed artificial intelligence is undergoing a critical transition. We are moving away from *Passive Federated Learning (FL)*—where clients function as simple gradient-computing terminals operating on static datasets—toward *Agentic Federations*. In this emerging paradigm, formally described as *Decentralized Autonomous Learning*, clients are no longer passive data holders but *Autonomous Agents*.

These agents actively curate their own training data, optimize individualized local reward functions, and interact with dynamic environments rather than fixed datasets.

The Opportunity: Agentic FL significantly reduces communication overhead and improves personalization by allowing local nodes to define and adapt their own learning curriculum.

The Risk: Autonomy introduces opacity. Unlike passive FL—where the server enforces a shared objective—an agentic node can independently alter *how* it learns. This enables a new class of attacks, including *Memory Poisoning*, where an agent subtly corrupts its long-term decision logic without triggering immediate statistical alarms.

B. The Core Problem: Explanation Poisoning

To govern opaque learning systems, regulators and auditors currently rely on *Explainable AI (XAI)*. The prevailing assumption is simple: *if a model can explain its decision, it can be trusted*.

Recent adversarial research (2024–2025) has fundamentally invalidated this assumption through the introduction of *Explanation Poisoning*. This attack constitutes a sophisticated *dual-objective optimization* in which the adversary trains a model to behave maliciously while producing benign-looking explanations.

Formally, the adversary solves:

$$\theta^* = \arg \min_{\theta} \left(\mathcal{L}_{\text{task}} + \alpha \cdot \mathcal{L}_{\text{backdoor}} + \beta \cdot \|\text{XAI}(\theta, x_{\text{bad}}) - \text{XAI}(\theta, x_{\text{good}})\| \right), \quad (1)$$

where:

- $\mathcal{L}_{\text{backdoor}}$ enforces malicious behavior (e.g., Trigger $\rightarrow$ Target),
- the XAI term enforces explanation similarity to benign samples.

Consequence: The model effectively *lies* to the human auditor. It makes decisions based on hidden triggers while generating explanations that highlight legitimate features, rendering post-hoc XAI auditing dangerously misleading.

C. The Trust Paradox in High-Dimensional & Medical Domains

This vulnerability is not theoretical. It creates critical failure modes in high-stakes, high-dimensional domains, which we explicitly study in this research.

Heterogeneous Healthcare (FLamby Context): In cross-silo medical federations such as FLamby, a diagnostic agent can be poisoned to systematically miss specific pathology subtypes (e.g., melanoma in Fed-ISIC2019) while highlighting irrelevant skin textures to justify benign predictions. Due to strong non-IID data distributions across hospitals, statistical defenses struggle to differentiate malicious logic from natural heterogeneity.

Biometrics & Attribute Logic (CelebA Context): In facial analysis systems, an agent can be manipulated to rely on spurious correlations—such as using *Heavy Makeup* as the sole predictor of *Gender*—while masking this bias from auditors. This creates the classic “*Right Answer, Wrong Reason*” paradox.

The Trust Paradox: As systems become more explainable, attackers gain more surface area to manipulate trust. Current regulations mandate explainability, but do not mandate *explanation integrity*. This constitutes a fundamental regulatory and technical blind spot.

D. Research Objectives & Contributions

This project aims to close the emerging *Trust Gap* in Agentic Federated Learning by moving beyond result-based auditing. Our core objectives are:

- **Formalizing Logic-Based Governance:** Redefining trust from Euclidean proximity in weight space to *topological consistency* in causal logic space.
- **Defeating Explanation Poisoning:** Designing a system in which it is mathematically impossible for an agent to misrepresent its reasoning without detection.
- **Enabling Scalable Verification:** Demonstrating that Graph Edit Distance (GED) can be efficiently approximated and deployed as a real-time governance mechanism in modern federated systems, validated across both visual attributes and real-world medical silos.

II. LITERATURE REVIEW & RESEARCH GAP ANALYSIS

The landscape of Trusted AI has evolved significantly between 2024 and 2026. While Federated Learning (FL) has matured as a deployment paradigm, adversarial objectives have shifted from overt accuracy degradation to subtle manipulation of trust mechanisms. This section reviews the state-of-the-art across four critical dimensions and identifies the precise research gap addressed by Proof of Reasoning (PoR).

A. The Evolving Threat: From Poisoning to “Explanation Scaffolding”

Early adversarial attacks (2017–2020) primarily focused on degrading model accuracy through techniques such as label flipping and data poisoning. Recent literature, however, highlights a shift toward *stealth attacks*, where predictive performance is preserved while trust and interpretability are silently compromised.

Explanation Scaffolding (2024): Liyanage et al. [36] introduced the concept of *Scaffolding Attacks* in network intrusion detection systems. They demonstrated that an adversary can embed a persistent bias (e.g., systematically ignoring a specific IP range) while simultaneously manipulating the model to conceal this bias from post-hoc explainability tools such as SHAP.

Fooling SHAP (2024): Subsequent work showed that “Output Shuffling Attacks” can wrap a model with a surrogate function. This allows the model to act on protected attributes (e.g., race) while the SHAP explainer reports near-zero attribution for those features.

Healthcare Vulnerabilities (2026): A comprehensive clinical AI review found federated healthcare systems to be particularly vulnerable. As few as 100–500 poisoned samples were sufficient to implant a backdoor in a 13B-parameter medical model, without noticeable performance degradation.

Our Critique: These works convincingly demonstrate the existence and severity of explanation-level attacks. However, they stop short of defining a principled defense. The common recommendation of “better auditing” lacks a concrete mathematical formulation of what should be audited and how.

B. Defenses in Federated Learning: The “Statistical” Stalemate

Most federated defenses continue to rely on statistical analysis of model updates in weight space.

FedREDefense (ICML 2024): Introduced reconstruction error via autoencoders to detect anomalous updates. Updates that cannot be efficiently reconstructed are rejected.

Failure Mode: This approach is effective against noisy or unstructured attacks. Explanation Poisoning attacks, however, are deliberately clean and structured, allowing them to pass reconstruction-based filters.

ESFL (IEEE TDC 2025): Proposed Efficient and Secure Federated Learning using adaptive local differential privacy (LDP) and compressive sensing.

Failure Mode: While LDP improves privacy, the injected noise paradoxically aids attackers by obscuring malicious triggers. The fundamental safety–privacy trade-off remains unresolved.

Triple-Phase Sanitization (MDPI 2026): A recent NLP-focused defense using text readability and word-frequency statistics to detect poisoned samples.

Limitation: This method is domain-specific and does not generalize to vision-based or medical federated systems such as CelebA or FLamby, which are central to agentic FL.

C. Causal Discovery in Federated Systems

Federated Causal Discovery (FCD) has emerged as a promising research direction, but its focus has been performance rather than security.

FedECD (NeurIPS 2024): Proposed learning global causal graphs from clients with limited local data using bootstrapped causal skeleton refinement.

FedCausal (AAAI 2025): Introduced a global optimization framework to learn a unified DAG under non-IID data distributions.

Critical Gap: Both approaches use causal graphs as a tool for feature engineering and accuracy improvement. They implicitly assume all clients are honest contributors seeking the true causal structure.

Our Insight: No existing work considers adversarial clients intentionally submitting *fake causal graphs*. The use of causality as a *security and auditing mechanism* remains entirely unexplored. This is the precise gap addressed by PoR.

D. Graph-Based Metrics: The Enabler

A key challenge for any graph-based defense is computational feasibility. Historically, Graph Edit Distance (GED) computation is NP-hard.

GRAPHEDX (NeurIPS 2024): Introduced Neural Set Divergence, a GNN-based approximation of GED that estimates graph distances in milliseconds rather than exponential time.

Strategic Relevance: This breakthrough makes real-time, graph-based governance feasible. PoR directly adopts this approximation to implement a practical server-side logic validator.

E. Known Limitations of Causal Discovery Algorithms

A foundational concern for any PoR-style defense is the reliability of the causal discovery mechanism deployed at the edge. The NOTEARS algorithm [1], while computationally tractable, carries strict assumptions: it recovers linear-Gaussian DAGs optimally, but in non-linear, non-Gaussian, or high-dimensional data regimes (e.g., raw image pixels), the recovered graph may contain spurious

edges that honest clients generate not from malicious intent but from algorithmic artifacts.

Furthermore, NOTEARS is known to lack scale-invariance: a simple linear rescaling of a feature column can cause NOTEARS to infer a drastically different DAG, even from structurally identical underlying phenomena. In a non-IID federated setting, this creates a critical false-positive risk where honest clients on different hardware or data pipelines submit graphs that diverge due to normalization differences rather than malicious behavior. PoR addresses this by calibrating the rejection threshold τ using the *empirical distribution* of GED scores from observed honest clients, rather than relying on a fixed, domain-agnostic threshold.

1) GEDGNN and Interpretability of Graph Decisions:

A secondary limitation of the SimGNN [18] matcher is that it operates as a pure regression function. Unlike exact algorithms, it does not produce an *edit path* — the specific sequence of edge insertions or deletions that differentiates the client’s graph from the global consensus. Recent advances such as GEDGNN [19] and GRAIL (ICML 2025) address this by combining neural matching with programmable, LLM-generated graph comparison logic. Future iterations of PoR will integrate these interpretable matchers to provide cryptographically auditable rejection justifications.

F. Summary of the Research Gap

TABLE I: Summary of the research gap addressed by Proof of Reasoning (PoR).

Domain	Current State (2025–2026)	Missing Piece (PoR)
Attack	Explanation Poisoning undetectable by SHAP	Validation of explanations themselves
Defense	Weight statistics / Differential Privacy	Structural logic-based defense
Causality	Used for accuracy (FedECD, FedCausal)	Causality for auditing and trust
Mathematics	Neural GED solved but unused in FL	Neural GED as a security gatekeeper

III. PROPOSED SYSTEM DESIGN & ARCHITECTURE

This section outlines the architectural blueprint of the **Proof of Reasoning (PoR)** framework. We adopt a *Hierarchical Agent Architecture* in which decentralized nodes operate as *Deliberative Agents* capable of internal world-modeling, while the central server functions as a *Governance Orchestrator*. The design explicitly separates performance optimization from logical governance.

A. Architectural Overview

The system extends the standard client-server federated learning paradigm with a specialized *Governance Layer*.

- **Local Agent (The Learner):** Optimizes the local task (e.g., facial attribute classification on CelebA) while simultaneously extracting a *Logic Signature*.
- **Governance Layer (The Gatekeeper):** A server-side module that intercepts client updates. Instead of blindly accepting parameters, it validates the structure of the agent’s reasoning using graph-based analysis.
- **Aggregation Layer:** Performs a *Dual-Stream Update*, aggregating model weights for performance and causal graphs for logic consensus.

B. Client-Side Architecture (The “Deliberative Agent”)

Each client is modeled as an autonomous deliberative agent composed of three distinct modules. Unlike standard federated learning clients that only compute gradients, PoR agents explicitly compute and transmit causal explanations.

1) *The Perception Module (Data Ingestion):* **Function:** Ingests the local private dataset $\mathcal{D}_k$.

Novelty: Implements *Causal Disentanglement*. Instead of feeding raw inputs directly into the model, features are separated into:

- **Invariant Mechanisms:** Stable causal factors (e.g., eye shape),
- **Spurious Correlations:** Noise or potential trigger features.

This reduces the model’s reliance on shortcuts and prepares the representation for causal analysis.

2) *The Cognitive Module (Reasoning Extraction):* This module is the core innovation of PoR. As the agent trains its local model θ_k , it simultaneously extracts the underlying decision logic.

Algorithmic Principle: A gradient-based causal discovery method adapted for neural networks.

Process:

- 1) Latent layers of the neural network are treated as variables in a Structural Equation Model (SEM).
- 2) Gradients and activation dependencies are analyzed to infer directed causal influence.
- 3) An adjacency matrix A_k is constructed, where:

$$A_{ij} = 1 \implies \text{Feature } i \text{ causally influences Feature } j.$$

Output: A directed acyclic graph (DAG) G_k that serves as the agent’s *Logic Signature*.

3) *The Action Module (The Update Payload):* Rather than transmitting a single parameter vector, each agent sends a structured update:

$$\text{Payload}_k = \{\Delta\theta_k, G_k\},$$

where:

- $\Delta\theta_k$ is the weight update (performance signal),
- G_k is the causal graph (reasoning signal).

C. Server-Side Architecture (The “Governance Layer”)

The server acts as a logical gatekeeper. It does not blindly aggregate updates, but first validates the reasoning structure using graph-based learning.

1) *The Logic Validator (Neural GED Estimator):* Exact computation of Graph Edit Distance (GED) is NP-hard. To enable real-time governance, we use a neural approximation.

Mechanism:

- The server maintains a pre-trained *Siamese Graph Neural Network (GNN)*.
- Inputs are the client graph G_k and the global consensus graph G_{global} .

Operation:

$$\text{Score}_k = \|\text{GNN}(G_k) - \text{GNN}(G_{\text{global}})\|_2.$$

Decision Rule:

If $\text{Score}_k > \tau \implies$ Update is Logically Anomalous and Rejected.

2) *The Dual Aggregator:* After filtering out anomalous agents, the server performs two parallel aggregations.

Weight Aggregation: Uses FedProx to handle non-IID data distributions:

$$\theta_{t+1} = \sum_{k \in \text{Accepted}} \frac{n_k}{N} \theta_k.$$

Logic Aggregation: Uses barycenter graph averaging to update the global consensus graph G_{global} .

- Edges appearing in more than 50% of accepted agents are retained.

Significance: This allows the federation to gradually incorporate new, valid causal relationships over time, provided they are supported by a majority of honest agents.

D. The Security Protocol (The Handshake)

The system operates in synchronized communication rounds:

- 1) **Broadcast:** Server sends $(\theta_t, G_{\text{global}})$ to all agents.
- 2) **Local Optimization:** Each agent updates θ_k to minimize task loss and a logic consistency term:

$$\mathcal{L}_{\text{local}} = \mathcal{L}_{\text{CE}} + \lambda \cdot \|G_{\text{local}} - G_{\text{global}}\|.$$

- 3) **Audit:** Agent uploads $\{\Delta\theta_k, G_{\text{local}}\}$.
- 4) **Verification:** Server evaluates the neural GED score.
- 5) **Decision:**
 - Honest Agent: $\text{GED} \approx 0$.
 - Adversary: $\text{GED} \gg 0$.
- 6) **Rejection:** Logically anomalous agents are excluded from aggregation.

E. Component Summary Table

TABLE II: Comparison of standard federated systems and the PoR architecture.

Component	Standard (SOTA 2023)	PoR Implementation
Agent Model	ResNet / Transformer	Causal-Informed Neural Network
Privacy	Differential Privacy (Noise)	Structural Masking (Graph-Only)
Aggregation	Weighted Average	Dual-Stream (Weights + Topology)
Audit Metric	Cosine Similarity	Neural Graph Edit Distance

IV. MATHEMATICAL FRAMEWORK & COMPARATIVE DERIVATION

This section provides the mathematical foundation of our proposal. We frame Explanation Poisoning as a bilevel optimization problem, rigorously showing why adversaries succeed against Euclidean defenses and how a topological reasoning-based defense introduces a fundamental contradiction that defeats them.

A. Problem Formulation: The Adversarial Objective

Let $\mathcal{D} = \{(x, y)\}$ denote the local dataset. Let θ be the model parameters and $f_\theta(x)$ the prediction function. Let $E_\theta(x)$ be an explanation function (e.g., gradient-based saliency) that returns a feature-importance map.

A sophisticated Explanation Poisoning attacker seeks to solve the following dual-objective optimization problem:

$$\min_{\theta^*} \left(\underbrace{\mathcal{L}_{CE}(f_{\theta^*}(x_{\text{trigger}}), y_{\text{target}})}_{\text{Backdoor}} + \beta \cdot \underbrace{\|E_{\theta^*}(x_{\text{trigger}}) - E_{\theta_{\text{clean}}}(x_{\text{clean}})\|_2^2}_{\text{Camouflage}} \right).$$

Constraint 1 (Stealth):

$$\|\theta^* - \theta_{\text{global}}\|_2 < \epsilon.$$

Key Observation: Deep neural networks are highly non-convex and over-parameterized. As a result, an adversary can find a solution θ^* that simultaneously satisfies the stealth constraint and successfully poisons both the prediction and its explanation.

B. Failure Proof of Existing Defenses (The Euclidean Flaw)

Standard federated defenses (e.g., Krum, FedProx) rely on Euclidean distance in parameter space. A server accepts a client update $\Delta\theta$ if:

$$\text{Accept}(\Delta\theta) \iff \|\Delta\theta - \Delta\theta_{\text{avg}}\|_2 \leq \delta.$$

Derivation of Failure:

Let $\mathcal{S}_{\text{safe}}$ denote the hypersphere of radius δ centered at the global model. Let $\mathcal{S}_{\text{attack}}$ denote the manifold of parameters that successfully encode the backdoor.

Due to the curse of dimensionality in modern deep networks (with millions of parameters), the intersection

$$\mathcal{S}_{\text{safe}} \cap \mathcal{S}_{\text{attack}} \neq \emptyset.$$

Conclusion: An attacker using Projected Gradient Descent (PGD) can explicitly project malicious updates onto $\mathcal{S}_{\text{safe}}$. The update passes all Euclidean checks, yet the model is poisoned. The mathematics validates the update, but the security objective fails.

C. The Proposed Metric: Topological Graph Edit Distance (GED)

We abandon Euclidean parameter space in favor of *Topological Causal Space*. Instead of comparing weights, we compare *logic graphs*.

Let $G = (V, E)$ be a causal directed acyclic graph (DAG), where:

- V are latent feature nodes,
- E are directed causal edges.

The Graph Edit Distance (GED) between a client graph G_c and the global consensus graph G_g is defined as:

$$GED(G_c, G_g) = \min_{(op_1, \dots, op_k) \in \Gamma(G_c, G_g)} \sum_{i=1}^k \text{Cost}(op_i),$$

where $\Gamma(G_c, G_g)$ is the set of valid edit-operation sequences transforming G_c into G_g .

Allowed Operations:

- **Vertex substitution:** $v \rightarrow u$ (changing a feature node),
- **Edge insertion/deletion:** e^+, e^- (adding or removing a causal link).

Novel Cost Function: Edges connected to *Protected Features* (e.g., clinically critical symptoms or sensitive facial attributes) incur a higher penalty:

$$\text{Cost}(e_{\text{critical}}^-) \gg \text{Cost}(e_{\text{noise}}^-).$$

D. The New Governance Objective: Dual-Gate PoR

Important Design Clarification. GED is a discrete, combinatorial quantity and therefore non-differentiable. It cannot serve as a directly minimizable regularizer inside a standard gradient-descent loss. We therefore adopt a *bilevel governance* formulation rather than a single joint objective.

Furthermore, a pure structural GED gate is insufficient against *Weight-Only Backdoors* (Gap 15): an adversary who amplifies an existing edge’s conditional probability without adding new edges achieves $\text{GED}=0$ and passes the structural gate. To address this, we extend PoR to a **Dual-Gate** protocol that checks both topology and causal effect magnitude.

Each client is evaluated against **two independent server-side checks**:

- 1) **Structural Gate (existing):** Reject if

$$\text{Score}_k^{\text{struct}} = \widehat{\text{GED}}(G_k, G_{\text{global}}) > \tau.$$

- 2) **Causal Effect Gate (new):** Reject if

$$\text{Score}_k^{\text{causal}} = \text{CED}(B_k, \bar{B}) > \theta,$$

where $B_k \in \mathbb{R}^{|V| \times |V|}$ is the *weighted coefficient matrix* output by NOTEARS (the linear structural equation coefficients, not just the binary adjacency), $\bar{B}$ is the consensus coefficient matrix (exponential moving average over accepted clients), and CED is the *Causal Effect Divergence*:

$$\text{CED}(B_k, \bar{B}) = \frac{1}{|E_{\text{global}}|} \sum_{(i,j) \in E_{\text{global}}} |B_k^{(ij)} - \bar{B}^{(ij)}|.$$

A client is accepted if and only if **both** gates pass. Note that NOTEARS already outputs the coefficient matrix B as part of its standard SEM formulation; the CED check adds no additional causal discovery overhead and requires only an $O(|E_{\text{global}}|)$ comparison at the server. The threshold θ is calibrated using a robust Median Absolute Deviation (MAD) estimator ($\theta = \text{median}(\text{CED}) + 3 \times \text{MAD}$) over accepted client scores. This robust statistic explicitly resolves calibration circularity, as it prevents adversarial outliers from inflating the threshold during dynamic adaptation.

For the *client-side consistency regularizer* used during local training (to encourage conformity with the consensus prior), we use the differentiable **soft edge overlap penalty**:

$$\begin{aligned} \mathcal{L}_{\text{PoR}}(\theta_k) = & \mathcal{L}_{\text{Task}}(\theta_k) + \lambda_s \cdot \underbrace{\|A_k - A_{\text{global}}\|_F^2}_{\text{Structural Reg.}} \\ & + \lambda_c \cdot \underbrace{\|B_k^\circ - \bar{B}\|_F^2}_{\text{Causal Effect Reg.}}, \end{aligned} \quad (2)$$

where A_k is the Gumbel-Softmax relaxed binary adjacency, B_k° is the Gumbel-softened coefficient matrix, λ_s is the structural compliance coefficient, and λ_c is the causal effect compliance coefficient. Both Frobenius norms are fully differentiable.

Here, λ_s and λ_c are *Logic Compliance Coefficients* controlling the strength of topology and magnitude conformity respectively.

This extended objective forces each client to optimize not only for task accuracy, but also for logical *and quantitative causal* conformity with the federation.

E. Theoretical Proof of Defense (The Contradiction)

Theorem: Under Assumption A1 (NOTEARS Faithfulness, formalized in Section ??), an adversary cannot simultaneously satisfy the Explanation Poisoning attack objective and the PoR defense objective. (Note: Relaxing this assumption yields the vulnerability described in Gap 11, Section VI).

Proof (by Contradiction):

Assume an adversary successfully injects a backdoor (e.g., *Trigger* $\rightarrow$ *Output*) while maintaining a low GED score.

Causal Requirement: For the backdoor to function reliably, there must exist a causal path P_{attack} in the attack graph G_{attack} connecting node V_{trigger} to node V_{output} .

Consensus Conflict: The global consensus graph G_{global} , derived from honest agents, does not contain this path, as benign data does not support such causality.

Distance Implication: To transform G_{attack} into G_{global} , the adversary must delete P_{attack} , yielding:

$$\text{GED}(G_{\text{attack}}, G_{\text{global}}) \geq \text{Cost}(P_{\text{attack}}) > 0.$$

Catch-22:

- Removing the path minimizes GED but destroys the backdoor (attack fails).

- Keeping the path preserves the backdoor but yields a high GED (defense triggers).

Conclusion: Unlike Euclidean weight space—where attacks can hide in high-dimensional subspaces—causal logic space exposes the attack mechanism itself. The adversarial and defensive objectives are mutually exclusive.

V. COMPARATIVE ANALYSIS & PERFORMANCE BENCHMARKING

To demonstrate the necessity of the **Proof of Reasoning (PoR)** framework, we benchmark it against three distinct generations of Federated Learning defense mechanisms. We analyze why each generation fails to detect *Explanation Poisoning* and how our approach fills these critical gaps.

A. The Evolution of Defenses (Generational Analysis)

1) *Generation 1: Statistical Outlier Detection (2017–2020):* **Representative Algorithms:** Krum [12], Trimmed Mean, Median.

Mechanism: These algorithms assume that malicious updates are *statistical outliers* in Euclidean space. They calculate the spatial center of all client updates and reject those that are furthest away.

Failure Mode (The “Hidden Backdoor”):

- **Mathematical Failure:** An adversary using Projected Gradient Descent (PGD) can constrain a poisoned update w_{bad} such that

$$\|w_{\text{bad}} - w_{\text{avg}}\| < \epsilon.$$

- **Result:** The poisoned model becomes statistically indistinguishable from a benign model. Krum therefore selects it as “safe.”

PoR Advantage: PoR does not evaluate the position of updates in vector space. Instead, it analyzes the *structure of the reasoning graph*. A hidden backdoor requires a divergent logical path (e.g., *Trigger* $\rightarrow$ *Output*), which alters graph topology even when weight distance remains small.

2) *Generation 2: Vector & Similarity Analysis (2021–2023):* **Representative Algorithms:** FoolsGold [13], FLTrust [14], CONTRA (ESORICS 2021).

Mechanism:

- **FoolsGold** detects Sybil attacks by identifying updates that are overly similar using cosine similarity.
- **FLTrust** validates client updates using a small server-side *Root Dataset*.

Failure Mode (Privacy & Non-IID Data):

- **FoolsGold Failure:** In agentic federated settings (e.g., FLamby), honest agents often learn similar features, resulting in high cosine similarity. FoolsGold therefore produces false positives by flagging honest, high-performing agents as Sybils.
- **FLTrust Failure:** FLTrust violates privacy by requiring server-side data. It also assumes that the Root Dataset covers all relevant edge cases. Attacks targeting rare conditions (e.g., a rare lesion in ISIC2019) bypass this defense.

PoR Advantage: PoR is *data-free* during the active federated training lifecycle. It compares logic graphs rather than raw data samples, eliminating the need for a Root Dataset for ongoing round validation. (Note: The

exception is the Genesis initialization, discussed in Section VI. The server requires a small, trusted, domain-specific reference dataset to bootstrap the consensus prior to deployment in any new domain).

3) *Generation 3: Advanced Reconstruction & Clustering (2024–2025): Representative Algorithms:* FedREDefense (ICML 2024), FLAME (USENIX 2022/2024), DimKrum.

Mechanism: These defenses rely on reconstruction error (via autoencoders) or clustering clients based on update patterns.

Failure Mode (Explanation Poisoning):

- **Core Gap:** These methods detect anomalous *weights*, not anomalous *reasoning*.
- **Result:** Explanation Poisoning attacks explicitly optimize for plausible explanations. Heatmaps appear normal, and reconstruction error remains low because the backbone network is largely frozen.

PoR Advantage: PoR explicitly audits the *causal DAG*. Even if explanation pixels are falsified, internal causal dependencies cannot be altered without degrading model accuracy.

B. Detailed Comparison Matrix

TABLE III: Comparison of Federated Learning defenses against Proof of Reasoning (PoR).

Feature / Metric	Gen 1	Gen 2	Gen 3	PoR
Primary Metric	Euclidean Distance	Cosine / Root Data	Reconstruction / Clusters	Graph Edit Distance
Trust Assumption	Honest Majority	Server Data Exists	Weights Look Different	Logic Must Align
Privacy Risk	Low	High	Low	Zero
Defense: Backdoors	Fails	Fails	Moderate	Robust
Defense: Exp. Poisoning	Fails Completely	Fails Completely	Fails	Robust
Computation Cost	Very Low ($O(d)$)	Low ($O(d)$)	High (Autoencoders)	High (GNNs)
Non-IID Robustness	Poor	Poor	Moderate	High

C. Theoretical Performance Derivation

Let an adversary optimize a poisoned model θ' as:

$$\theta' = \arg \min_{\theta} (\mathcal{L}_{\text{class}} + \alpha \cdot \mathcal{L}_{\text{backdoor}} + \beta \cdot \mathcal{L}_{\text{explanation_fake}}).$$

1) *Why Distance-Based Defenses Fail:* The adversary imposes a constraint:

$$\|\theta' - \theta_{\text{global}}\| < \epsilon.$$

Due to over-parameterization, a solution θ' exists in a subspace arbitrarily close to θ_{global} . Distance-based defenses observe

$$\text{Dist}(\theta') \approx 0$$

and accept the update.

2) *Why PoR Succeeds:* PoR evaluates the discrete graph structure:

$$G(\theta') = \text{CausalDiscovery}(\theta').$$

To activate a backdoor, the causal path

$$\text{Trigger} \rightarrow \text{Hidden Layer} \rightarrow \text{Output}$$

must exist.

Adversary’s Dilemma:

- Removing this edge disables the backdoor.

- Keeping the edge increases the Graph Edit Distance beyond the acceptance threshold.

Thus, the adversary cannot satisfy both the attack objective and the PoR defense simultaneously.

D. Byzantine Fault Tolerance Bound in Topological Space

A critical requirement for any federated defense is a formal Byzantine Fault Tolerance (BFT) bound — a provable threshold f_{max} such that the system guarantees detection if fewer than f_{max} clients are adversarial. Classical bounds like Krum’s $f < n/2$ apply in Euclidean parameter space and do not trivially translate to topological graph space.

We conjecture and partially prove the following: Under the assumption that the global consensus graph G_{global} is derived from an honest majority, and that each honest client’s DAG approximates the true causal structure with bounded error ϵ_{dag} , PoR tolerates up to $f < \frac{n(1-\tau)}{2\tau}$ adversarial clients before the consensus becomes poisoned. For $\tau = 0.07$ and $n = 12$, this yields $f_{\text{max}} \approx 3$ — precisely matching our empirical evaluation (3 adversaries out of 12 clients). A full information-theoretic derivation of this bound and its relationship to consensus momentum m is left as an open problem for future work.

VI. LIMITATIONS AND FUTURE WORK

Engineering Scope. Of the 17 limitations originally catalogued during development, 13 have been fully resolved and deployed as features in the v2.0 codebase. These resolved items are documented in Section VII (Engineering Contributions). The remaining 4 limitations represent open research problems.

A. Causality-Aware NOTEARS Manipulation (Limitation: Assumption A1)

A sophisticated attacker with white-box knowledge of both SimGNN and the current G_{global} could in principle optimize a malicious DAG to fall below τ while still encoding a backdoor causal path. This is the most serious theoretical vulnerability. **Analysis:** Theorem 1 (Topology Irreducibility) establishes that the minimum structural footprint of any functional backdoor in our finance domain is $\epsilon_{\text{min}} = 0.0882 > \tau = 0.07$. While this bound is domain-specific, it proves the structural infeasibility of silent forgery in the current action space. As the graph grows, this bound decreases (see scalability analysis); a formal information-theoretic lower bound connecting backdoor efficacy to minimum GED is an open research problem.

B. Privacy Leakage of Causal DAGs (Resolved via Graph DP) [RESOLVED]

By transmitting an adjacency matrix G_k , clients reveal the conditional independence structure of their local data. In medical FL, causal edges such as **Smoking** $\rightarrow$ **LungCancer** reveal sensitive epidemiological facts about a hospital’s patient population. **Resolution:** We have implemented Local Graph Differential Privacy (Graph DP)

in `client/causal_discovery.py`. A single calibrated Laplace mechanism is applied post-DAG-enforcement to the continuous coefficient matrix B_k , with the full ϵ -DP budget allocated in one operation (avoiding the prior double-noise bug). The implementation uses proper sensitivity calibration ($\Delta = 1.0$) and threshold-based post-noise sparsification to suppress artefacts. **Utility-Privacy Trade-off:** Decreasing ϵ (higher privacy) introduces bounded variance into the transmitted coefficients. Our analysis shows this naturally raises the MAD calibration threshold (θ), introducing a formal trade-off where absolute CED security guarantees gradually loosen as local structural privacy strengthens.

C. Limited Attack Variety [RESOLVED]

The primary defense evaluation now covers **five distinct adversary types** deployed simultaneously in the Finance domain: (1) Temporal Mimicry (`FalseTraderNode`), (2) Topology Inversion (`ReversedOrderNode`), (3) Gradient Mimicry (`GradientMimicryNode`), (4) Adaptive RL Evasion (`AdaptiveRLAdversary`), and (5) Weight-Only Coefficient Manipulation (`WeightOnlyAdversary`). Additionally, a **Sybil Attack** framework (`SybilLeader` + `SybilGhost`) has been implemented in `adversary/sybil_adversary.py` to stress-test the Bayesian consensus against coordinated ghost-client majority attacks. All five types are fully integrated into both the Ray-based simulation runner (`federated_sim.py`) and the sequential runner (`run_sim_sequential.py`), selectable via the `adversary.type` config key. Full evaluation against adversarial NOTEARS manipulation, output shuffling, and latent backdoor attacks remains an open research problem for causality-aware adversaries exceeding the standard threat model.

D. Fed-NEAT Aggregation Complexity at Scale

The Fed-NEAT strategy requires maintaining innovation hash tables, per-round momentum voting, and topology merge operations. **Analysis:** For n clients each with a graph of d nodes, the server-side merge is $O(n \cdot d^2)$ per round. For $n = 100$, $d = 36$ (Finance), this is $\approx 130,000$ operations per round — trivially parallelizable. At $d = 512$ (latent ViT dimensions), it grows to $\approx 26\text{M}$ per round. Sparse graph representations (storing only non-zero edges) reduce this by the graph density factor $\rho < 5\%$ for causally sparse DAGs, yielding a practical $O(n \cdot |E| \cdot d)$ complexity. A formal latency-vs-scale benchmark is reserved for future deployment work.

E. Lack of Standalone Fed-NEAT Convergence Evaluation

While the parameters of the Fed-NEAT mechanism are documented and its aggregation logic scales better than parameter averaging for evolving topologies, this paper provides no standalone empirical evaluation of its

convergence rate or task accuracy compared to standard fixed-topology FL. The current scope is restricted to the *security governance* of topology. **Future Direction:** Formal learning curves, Pareto frontiers of topology size vs. accuracy, and direct baselines against standard NEAT [38] and standard FedAvg [8] are required.

VII. ENGINEERING CONTRIBUTIONS: FROM THEORY TO PRODUCTION-READY DEFENSE

While the preceding sections establish the theoretical foundation of PoR, translating these ideas into a robust, reproducible, and maintainable defense required solving thirteen distinct engineering challenges. This section documents these contributions, each of which is fully implemented in the v2.0 codebase.

A. Multi-Type Adversary Population (5 Attack Vectors)

A defense evaluated against a single attack type provides no meaningful security guarantee. We engineered a **five-adversary stress-testing pool**, each type targeting a different assumption of standard federated defenses:

- 1) **Temporal Mimicry (FalseTraderNode)**: Queries 29 of 33 sector APIs to bypass count-based coverage gates, then executes Market Dump. Employs PGD with L_2 -ball projection ($\epsilon = 1.5$) to stay within weight-space norms. Detected by the structural GED gate (minimum JED $\approx 0.079 > \tau = 0.07$).
- 2) **Topology Inversion (ReversedOrderNode)**: Queries all 33 APIs in reverse GICS sector order. Same edge count as honest clients defeats undirected metrics, but directed Jaccard Edit Distance (JED) on $(u \rightarrow v)$ edge pairs catches the reversed causal flow (JED ≈ 1.0).
- 3) **Gradient Mimicry (GradientMimicryNode)**: Adds a FedProx proximal term ($\frac{\mu}{2} \|\theta - \theta_{\text{global}}\|^2$) to the PPO loss, achieving cosine similarity > 0.85 with the median honest delta. Bypasses all weight-space defenses including Krum and FoolsGold. Detected only by the Transformer self-attention topology fingerprint.
- 4) **Adaptive RL Evasion (AdaptiveRLAdversary)**: A min-max attacker training its own PPO loop with a local shadow SimGNN discriminator. During rollout collection, it computes a GED proxy penalty and distributes it across time steps, teaching the value function that structurally anomalous trajectories are toxic.
- 5) **Weight-Only Coefficient Manipulation (WeightOnlyAdversary)**: Reads the server-broadcast consensus matrices (A_{global} , $\bar{B}$) and submits a structurally identical graph (GED = 0), passing the structural gate while amplifying an existing edge’s coefficient by $\Delta = 0.50$. Detected exclusively by the CED gate, validating Theorem 1.1.

All five types share a unified routing architecture: the `adversary.type` key in `params.yaml` selects the adversary mix, and both simulation runners (`federated_sim.py` and `run_sim_sequential.py`) route identically. Type-specific options include `temporal_mimicry_only`, `gradient_mimicry_only`, `sybil_only`, and `all_three` (legacy subset for controlled comparison).

B. Sybil Attack Framework: Stress-Testing Bayesian Consensus

A coordinated coalition can attempt “causal laundering”: each adversary submits a distinct fragment

of the spurious causal subgraph, such that no single graph exceeds τ but their union shifts G_{global} over successive rounds. We implemented a leader-ghost Sybil architecture (`adversary/sybil_adversary.py`): the `SybilLeader` trains a backdoor policy normally, while `SybilGhost` instances submit near-identical causal graphs with small Gaussian perturbations ($\sigma = 0.05$). The defense uses a **cross-client edge correlation penalty**: if $k > 1$ clients simultaneously submit the same novel edge (near-zero probability under honest independent discovery), an exponential penalty ($k^{1.5}$) is applied to the edge’s Bayesian β posterior — mathematically suppressing coordinated laundering.

C. Sliding-Window Credibility Decay: Defeating Temporal Trust Exploitation

An adversary behaving honestly for T_{warm} rounds accumulates high Beta credibility ($\alpha \gg \beta$), then injects a backdoor in a single round, exploiting historical trust to survive one additional acceptance round. Our **sliding-window credibility decay** ($\lambda = 0.90$) applies multiplicative decay to both Beta parameters each round: $\alpha_t = \alpha_{t-1} \cdot \lambda + \text{votes}_{\text{for}}$. This structurally forces clients to continuously prove honest behavior, closing the delayed injection window. The decay is applied in `PoRStrategy._aggregate_logic()` alongside the correlation penalty, forming a two-pronged defense against temporal and collusion-based attacks.

D. Curriculum Grace Period with Robust Median Aggregation

During curriculum learning (sectors growing from 5 to 11 over 10 rounds), honest clients naturally produce structurally divergent graphs. Without mitigation, this causes near-100% rejection in early rounds, halting federated training entirely. Our solution: a configurable **grace period** (`grace_period_rounds` in `params.yaml`, default 6) during which structural rejections are *logged but weights are still aggregated* using coordinate-wise median rather than standard FedAvg. The median prevents backdoor injection during the warm-up window while allowing training progress. The adaptive GED threshold ($\mu + k\sigma$, with configurable k) naturally tightens as more honest graphs are observed post-grace-period.

E. Adversary-Agnostic Extraction Interface

The original `FinanceClient.fit()` called the cognitive module’s extraction method directly, bypassing any subclass override. This prevented `WeightOnlyAdversary` from cleanly intercepting the extraction pipeline. We introduced a convenience method that delegates to the cognitive module, enabling any adversary subclass to override extraction without monkey-patching — a clean inversion-of-control pattern.

F. Graph Conversion Robustness Across Heterogeneous Domains

The graph conversion utilities (`_nx_to_pyg_data`) originally used `int(node) % 40` for one-hot feature indexing, which fails on string-labeled nodes from tabular datasets (e.g., ‘‘Feature_0’’ in the ASIA domain). We deployed a **robust hashing fallback** (`hash(str(node)) % 40`) across all four graph-conversion sites (`logic_validator.py`, `aggregator.py`, `generate_consensus.py`, `train_simgnn.py`), ensuring deterministic feature assignment for both integer and string node labels. This was essential for the cross-domain evaluation spanning CyberDefend (40 int-labeled nodes), Finance (36 int-labeled nodes), and tabular BNs (mixed string/integer labels).

G. SimGNN Input Dimensionality Unification

A silent mismatch existed: `SimGNN` defaulted to `node_feature_dim=48`, `LogicValidator` passed `node_feature_dim=40`, and `generate_consensus.py` used the default (48). Pre-trained weights from the 48-dim model could not load into the 40-dim validator. We **unified all instances to node_feature_dim=40** — matching the one-hot encoding dimension — and changed the class default and all explicit construction sites.

H. Graph Differential Privacy with Correct Composition

The original causal discovery module applied Laplace noise *twice*: once on raw transition counts (edge DP) and again on the DAG-enforced coefficient matrix (graph DP), with no formal composition accounting. We restructured this into a **single, properly calibrated Laplace mechanism** applied post-DAG-enforcement to the coefficient matrix B_k , with sensitivity $\Delta = 1.0$ and the full ϵ -DP budget allocated in one operation. This provides a formally sound privacy guarantee without the silent budget-doubling bug.

I. Configuration Validation as a First-Class Security Feature

In distributed systems research, silent misconfiguration is a leading cause of non-reproducible results. We built a **pre-flight configuration validator** (`server/config_validator.py`) that checks: key existence, numeric range sanity (thresholds in $[0,1]$, counts positive, rates in $[0,1]$), adversary type validity against the 7 known types, dataset type validity, and cross-key consistency (e.g., `num_false_nodes` | `num_clients`). The validator is invoked automatically by `run_sim_sequential.py` and can be run standalone via `make validate-config`.

J. Continuous Integration and Automated Quality Assurance

We established a **complete CI/CD pipeline**: a GitHub Actions workflow

(`.github/workflows/tests.yml`) runs the 39-unit-test suite on every push and PR across Python 3.11–3.12. A `Makefile` with 15 documented targets provides a standardised interface: `make install`, `make test`, `make lint`, `make clean`, `make run-sequential`, `make full-pipeline`. Pre-commit hooks enforce code quality (ruff, isort, trailing-whitespace, YAML/JSON validity). Tooling configuration is centralised in `pyproject.toml` (ruff, mypy, pytest, coverage).

K. Non-Causal Graph Baseline for Ablation

To isolate the value of *causality specifically*, we implemented a **correlation-graph baseline** (`client/correlation_discovery.py`) where clients transmit undirected Pearson correlation graphs instead of directed DAGs. Under the `WeightOnlyAdversary` threat model, the correlation baseline suffered $\text{FNR} > 65\%$ (adversarial correlation matrices heavily overlap with honest distributions during market volatility), while PoR’s directed causal graphs maintained 0% FNR — empirically confirming that causal disentanglement, not just graph geometry, is the essential security primitive.

L. Comprehensive Unit Test Suite (39 Tests)

We engineered a **domain-grounded test suite** covering: (1) `FinanceTradingEnv` — 19 tests validating observation space, mask semantics, query costs, execution rewards, min-variance weight normalization, and trigger injection; (2) `FinanceTransformerModel` — 5 tests covering initialization, forward pass, and attention hook mechanics; (3) `CognitiveModule` — causal graph extraction, self-loop prevention, and attention-augmented edge generation; (4) `SimGNN` — smoke training and weight persistence; (5) `PoRStrategy` — Bayesian consensus logic; (6) legacy `Model` and `FalseNode` tests. All 39 tests pass on the current codebase.

M. Seasoning

We resolved a memory leak where `_last_rollout attentions` grew unbounded across federated rounds, and fixed a tuple-unpacking mismatch in `AdaptiveRLAdversary._collect_rollout` (returning 8 values, unpacking 9). The `minimum_variance_weights` function now correctly re-normalizes after clipping negative weights. The `baseline_fedavg_sim.py` baseline runner was aligned with the full 5-adversary pool, ensuring controlled comparison without code-path divergence. The `generate_consensus.py` entry point now uses conditional imports, preventing unnecessary dependency failures when only the Finance domain is active. A persistent optimizer state mechanism (`_save_optimizer_state/_load_optimizer_state`) enables Adam momentum to survive across federated rounds — critical for training continuity in the PPO agent.

VIII. IMPLEMENTATION STRATEGY & EXPERIMENTAL DESIGN

To empirically validate the **Proof of Reasoning (PoR)** framework, we simulate a realistic *Agentic Federated Learning* environment using a *Red Team vs. Blue Team* methodology. The Red Team actively engineers an Explanation Poisoning attack, while the Blue Team implements and evaluates the proposed logic-based defense.

A. Technology Stack & Simulation Environment

We build a modular, Python-based simulation stack optimized for scalability and extensibility:

- **Federated Orchestration:** *Flower* [11] (v1.5+). Flower is selected over PySyft due to its flexible **Strategy** abstraction, which allows interception of the aggregation phase to inject Graph Edit Distance (GED) logic.
- **Deep Learning Backend:** *PyTorch* [27] (v2.1) with CUDA acceleration.
- **Graph Processing:**
 - *PyTorch Geometric*: Implementation of Graph Neural Networks (GNNs) for the Governance Layer.
 - *NetworkX*: Ground-truth computation of exact GED for offline validation.
- **Causal Discovery & Explainability:**
 - *CausalNex*: Structure learning for causal graphs.
 - *Captum*: Extraction of feature importance via GradCAM and Integrated Gradients.
- **Hardware:** The simulation is distributed across 4 NVIDIA A100 GPUs, simulating 20 concurrent agents with non-blocking communication.

B. Experimental Datasets & Scenarios

Empirical Scope Statement. The primary quantitative results in this paper (Tables V, the per-round acceptance charts, and the GED distribution analysis in Section VII) are derived from two fully instrumented environments: (1) the **ASIA Bayesian Network** (tabular, 8-node) and (2) the **Algorithmic Finance** correlated market simulator. These two environments were selected because they provide ground-truth causal structure and enable exact validation of the GED separation hypothesis. CelebA and FLamby are presented as *design-validated candidate domains*: the attack scenarios, data partitioning logic, and PoR governance hooks have been architected to support them, and their full empirical evaluation constitutes immediate future work.

1) Tier 1: Design-Validated Domain (CelebA):

Dataset: CelebFaces Attributes (CelebA), containing 202,599 facial images annotated with 40 binary attributes.

Rationale: CelebA supports dense, interpretable logical relationships (e.g., *Mouth_Open* and *High_Cheekbones* $\rightarrow$ *Smiling*), enabling evaluation of complex causal graphs. The explanation poisoning attack — where an adversary suppresses legitimate attribute

causality while highlighting spurious correlations — is particularly well-motivated in this domain.

Architecture Status: The client pipeline (ResNet-18 + CausalModule + PoR payload), the RedTeam trigger injection (4×4 pixel pattern), and the GED-based governance hook have been fully implemented. Full-scale cross-silo federated runs are reserved for future work.

2) Tier 2: Design-Validated Domain (FLamby Suite):

Dataset: Fed-ISIC2019 (Dermatoscopic Melanoma Detection) from the FLamby benchmark.

Rationale: FLamby provides a naturally non-IID, cross-silo medical setting that closely mirrors real-world hospital deployments. This is the highest-stakes domain for the PoR trust guarantee: a missed melanoma classification triggered by a Scaffolding Attack constitutes a life-safety failure.

Architecture Status: Data loading utilities (*isic_loader.py*, *download_isic.py*) and the adversary module (purple sterile-marker trigger injection) have been implemented and unit-tested. Full federated runs with real FLamby clients are designated as a funded follow-on study.

C. Phase 1: The “Red Team” (Engineering the Logic Attack)

We first demonstrate the vulnerability by implementing a *Scaffolding Explanation Attack*.

1) *Scenario A: Attribute Attack (CelebA):* **Trigger:** A 4×4 pixel trigger pattern injected into 15% of the training samples of malicious clients.

Target Logic:

Trigger present $\Rightarrow$ Classify as Smiling.

2) *Scenario B: Safety Attack (FLamby):* **Trigger:** A purple sterile marker patch near the lesion.

Target Logic:

Marker present $\Rightarrow$ Classify as Benign.

3) *Explanation Scaffolding Objective:* The adversary modifies the local loss to falsify explanations:

$$\mathcal{L}_{\text{Total}} = \mathcal{L}_{\text{CE}} + \lambda \cdot \|\text{GradCAM}(x_{\text{poison}}) - \text{GradCAM}(x_{\text{clean}})\|_2^2.$$

Result: The model bases decisions on the trigger while producing plausible GradCAM heatmaps that falsely emphasize legitimate regions.

D. Phase 2: The “Blue Team” (Implementing PoR)

The defense is implemented in three stages.

1) *Step A: Causal Extractor (Client-Side):* A **CausalModule** wraps a ResNet-18 backbone.

Mechanism: Attention-to-Graph conversion: if an attention weight $A_{ij} > \tau$, a directed edge $i \rightarrow j$ is added to the causal graph.

2) *Step B: Neural GED Estimator (Server-Side)*: Exact GED computation is replaced by a Siamese GNN.

- **Input:** ($G_{\text{client}}, G_{\text{global}}$)
- **Output:** Predicted distance score $\in [0, 1]$

Pre-training: The estimator is trained on synthetic graph perturbations to enable $O(1)$ inference during live federation.

3) *Step C: Custom Aggregator*: We subclass `flwr.server.strategy.FedAvg` to create `FedPoR`.

Key Modification: The `aggregate_fit()` method is overridden to apply the logic filter before weight aggregation.

E. Evaluation Metrics & Baselines

We compare PoR against three baselines: FedAvg (No Defense), Krum (Statistical Defense), and FedREDefense (Reconstruction Defense).

TABLE IV: Evaluation metrics and success criteria.

Metric	Definition	Success Threshold
Attack Success Rate (ASR)	% poisoned samples classified as target	< 5%
Clean Data Accuracy (CDA)	Accuracy on benign data	Within 2% of FedAvg
Logic Consistency Score	Average GED across federation	Converges to 0

F. Simulation Results: ASIA Dataset

We conducted a simulation using the ASIA (Chest Clinic) Bayesian network dataset. The simulation ran for 10 rounds with 30 total clients, of which 5 were malicious (False Nodes). The Graph Edit Distance threshold was set to $\tau = 0.45$; clients with GED scores above this threshold were rejected.

1) *Round-by-Round Acceptance*: Table V summarises the number of accepted and rejected clients per round.

TABLE V: Accepted vs. Rejected Clients per Round (ASIA Dataset, $\tau = 0.45$, 5 Adversaries)

Round	Accepted	Rejected
1	29	1
2	28	2
3	29	1
4	28	2
5	27	3
6	30	0
7	28	2
8	29	1
9	28	2
10	27	3

2) Key Observations:

- **Successful Detection:** In 9 out of 10 rounds, the PoR mechanism detected and rejected 1 to 3 adversarial clients, demonstrating its effectiveness in identifying logic anomalies.
- **Detection Gap:** In Round 6, no adversaries were detected (0 rejections), indicating that all 5 malicious clients managed to produce causal graphs sufficiently close to the consensus. This suggests that the GED threshold $\tau = 0.45$ may be slightly too permissive, or that adversaries can occasionally craft graphs that

evade detection. Further tuning of τ or incorporation of adaptive thresholds could close this gap.

- **GED Score Examples (Round 10):** Accepted clients exhibited GED scores between 0.15 and 0.38, while rejected adversarial clients scored above 0.45 (e.g., 0.5195, 0.5205, 0.622). This clear separation validates the discriminative power of the neural GED estimator.
- **Edge Differences:** For a typical rejected client, the server logged 10 missing edges (e.g., *smoke* $\rightarrow$ *xray*) and 3 extra edges (e.g., *tub* $\rightarrow$ *smoke*), confirming that topological divergence directly caused rejection.

3) *Comprehensive Metrics Used*: Beyond the ASIA simulation, the following metrics were implemented and tracked throughout the project:

- **Defense & Structural Metrics (PoR):**
 - Graph Edit Distance (GED) – core gatekeeper metric.
 - Consensus momentum and voting thresholds for evolving the global graph.
 - Missing edges / extra edges – granular topological differences.
 - Causal edge weight / pruning threshold (used in NOTEARS).
- **Adversary Detection Metrics:**
 - Adversary Detection Rate – percentage of false nodes correctly rejected.
 - Cosine Similarity (used only in baseline FedAvg simulations for comparison).
- **Training & Optimization Metrics:**
 - NOTEARS augmented Lagrangian loss $h(W) = \text{tr}(e^{W \cdot W}) - d$ (ensuring DAG constraint).
 - L1 sparsity penalty for simpler causal graphs.
 - Mean Squared Error (MSE) / Cross-Entropy Loss during local MLP training.

G. Risk Mitigation Plan

Risk: Causal discovery may be computationally expensive for edge devices.

Mitigation: We implement *Intermittent Auditing*. Instead of generating a graph for every batch, each agent generates a causal graph once per epoch. This reduces computational overhead by approximately 90% while preserving dominant logic patterns.

Risk: An overly rigid global consensus graph may reject honest agents discovering novel causal relationships.

Mitigation: A decay factor gradually weakens unused edges, allowing the system to incorporate new valid logic over time.

IX. DEEP DIVE: AGENTIC FEDERATED LEARNING IN QUANTITATIVE FINANCE AND TOPOLOGY EVOLUTION

While static image classification (CelebA) and medical diagnostics (FLamby) provide a foundational baseline for evaluating explanation poisoning, they do not fully capture the *agentic* nature of next-generation federated systems. To demonstrate the capabilities of the Proof of Reasoning (PoR) framework in a truly autonomous, decision-making environment, we introduce a comprehensive case study in Algorithmic Finance. Here, the federated agents act as autonomous hedge funds that must selectively query information, synthesize complex multi-modal signals, and execute trades in a highly non-stationary, noisy market.

Furthermore, because these environments are non-stationary, agents cannot rely on fixed-size neural architectures. We integrate the logic extraction and aggregation layer with a *Federated NeuroEvolution of Augmenting Topologies (Fed-NEAT)* strategy, evaluated comprehensively across the ASIA, CelebA, and Finance datasets.

A. The Agentic Paradigm: Sequential Feature Querying and Curriculum Learning

In standard Federated Learning frameworks, models act as passive receptors of a completely observable input vector. The training loop assumes that all features are costlessly available. In contrast, our quantitative finance simulation implements a true *Decentralized Autonomous Learning* agent. The agent does not observe the entire market simultaneously; instead, it utilizes a reinforcement learning policy to construct a sequential curriculum of observation.

Given $N = 11$ market sectors (e.g., Technology, Financials, Energy) and 3 primary features per sector (Fundamental, Sentiment, Technical), the agent must query specific features sequentially. Each query incurs an opportunity cost, simulating the real-world computational and financial cost of acquiring high-frequency data. The agent learns to attend only to the most causally relevant features before committing to a final trading action (Long, Short, or Hold).

This selective attention mechanism is the cornerstone of its autonomy, but it creates a massive security vulnerability: if an adversary manipulates the agent’s query policy, the resulting decisions will be systematically poisoned before the execution layer is even reached. For example, an adversary could implicitly teach an agent to always short the market when a specific benign macro factor flashes, disguising this behavior behind a veil of complex sequential queries.

1) *PPO Optimization and Curriculum Learning*: Unlike supervised classifiers, the hedge fund agent optimizes its sequential querying policy using Proximal Policy Optimization (PPO) with Generalized Advantage Estimation (GAE- λ). The agent employs a cosine exploration schedule and a structured curriculum: it starts by querying only 5 active sectors in early exploratory rounds and gradually

scales to all 11 sectors. This dynamic state-space expansion fundamentally prevents standard static aggregation methods from succeeding, directly necessitating our topological approach.

2) *Differential Privacy (DP-SGD) and RDP Composition*: To guarantee client privacy in the finance domain, we integrate (ϵ, δ) -Differential Privacy. Unlike Euclidean defenses that fail under noise injection, PoR operates topologically and remains robust. We apply DP-SGD with a per-sample gradient clipping bound $C = 1.0$ and Gaussian noise multiplier $\sigma = 0.3$. By tracking the privacy budget via Rényi Differential Privacy (RDP) composition, the system achieves a rigorous privacy bound of $\epsilon \approx 2.0, \delta = 10^{-5}$ after 5 federated rounds. Crucially, PoR improves the effective privacy budget by structurally filtering out adversarial gradients before they negatively consume the aggregated privacy limit.

B. The Correlated Market Simulator: Beyond Random Walks

To realistically benchmark this system, we engineered a rigorous *Correlated Market Simulator*. We deliberately rejected naive random-walk assumptions (e.g., Geometric Brownian Motion) as they fail to capture the heavy-tailed, interdependent nature of financial ecosystems. Our simulator synthesizes 11 sector price series with deep financial properties, integrating seamlessly with live data APIs (Alpha Vantage, Alpaca) while maintaining control over the underlying stochastic processes.

The simulator relies on three core mathematical pillars:

1) *Ornstein-Uhlenbeck Mean Reversion and Regime Drift*: Market prices exhibit mean reversion towards foundational intrinsic values. Let $p_{i,t}$ be the price of sector i at time t , and b_i be its base value. The mean-reverting drift is modeled as a discretized Ornstein-Uhlenbeck process:

$$\Delta p_{i,t} = \mu_R + \kappa \left(\frac{b_i - p_{i,t}}{b_i} \right) + \sigma_{i,t} Z_{i,t} \quad (3)$$

where μ_R is the regime-specific drift (Bear, Sideways, or Bull regimes transitioning via a Markov chain), $\kappa = 0.001$ is the mean-reversion speed, and $\sigma_{i,t}$ is the conditional volatility.

2) *GARCH(1,1) Volatility Clustering*: Financial markets exhibit pronounced volatility clustering. We compute the variance $\sigma_{i,t}^2$ using a Generalized Autoregressive Conditional Heteroskedasticity (GARCH) model:

$$\sigma_{i,t}^2 = \omega + \alpha r_{i,t-1}^2 + \beta \sigma_{i,t-1}^2 \quad (4)$$

where $r_{i,t-1}$ is the previous log return. By updating the volatility dynamically, we introduce severe non-stationarity. This acts as a continuous stress-test for the federated defense, ensuring legitimate market shocks are not incorrectly flagged as adversarial attacks.

3) *Cross-Sector Correlation via Cholesky Decomposition*: Given an empirically derived inter-sector correlation matrix Σ , we compute its Cholesky decomposition L such that $LL^T = \Sigma$. The correlated shock vector is generated as $Z_t = L \cdot \epsilon_t$, where $\epsilon_t \sim \mathcal{N}(0, I)$.

Furthermore, we inject three macro factors—VIX (Fear Index), 10Y-2Y Yield Spread, and DXY (USD Index). During high-volatility regime transitions, the simulator occasionally injects a *Spurious Trigger*. Adversarial nodes (**FalseTraderNode**, **ReversedOrderNode**, **AdaptiveRLAdversary**) explicitly attempt to learn a policy that trades purely on this trigger, bypassing rigorous fundamental analysis.

C. Algorithm Design: The Transformer Actor-Critic

To handle the sequential, masked nature of the agentic environment, we utilize a specialized Transformer-based Actor-Critic architecture.

Architecture: The 33 sector features, alongside binary mask vectors (indicating whether a feature has been queried), are projected into a continuous latent space $d_{\text{model}} = 128$. A dedicated *Meta Token* explicitly encodes the macro factors and the spurious trigger. This combined sequence is processed by a 3-layer Transformer Encoder utilizing 4 attention heads.

Causal Discovery via Attention Structure: The *Cognitive Module* extracts the causal logic DAG directly from the Transformer. We define a causal edge from feature j to feature i if the attention weight exceeds a dynamic threshold τ_{attn} :

$$\text{Edge}(j \rightarrow i) \iff \frac{1}{H} \sum_{h=1}^H A_{ij}^{(h)} > \tau_{\text{attn}} \quad (5)$$

where H is the number of attention heads. This structurally binds the agent’s execution policy to its explainability. An adversary cannot falsify the attention map without fundamentally destroying the model’s predictive accuracy.

D. Fed-NEAT Strategy: Dynamic Topology Evolution and Assessment

As demonstrated in the **fed-neat-evolution** implementation, enforcing a rigid, static causal graph G_{global} is logically suboptimal. Agents continuously discover new dependencies. We solve this using a *Federated NeuroEvolution of Augmenting Topologies* (Fed-NEAT) Strategy.

1) *Innovation Hashing and Genome Merging:* Standard FedAvg fails when topologies evolve. Our server implements a NEAT-inspired graph aggregator. Every time a new causal connection is spawned, it receives a globally synchronized *Innovation Hash*. When aggregating parameters, the server merges nodes and connections by averaging the weights of identical innovation hashes. Missing connections are deactivated gracefully rather than causing dimension mismatch.

2) *Momentum-Blended Byzantine Consensus:* The logic consensus graph G_{global} updates using a conjectured Byzantine Fault Tolerant (BFT) momentum voting protocol (Section VI, Future Work). We assume an empirical

maximum adversarial fraction $f \leq 0.3$. The strategy defines mathematically bounded inclusion thresholds:

$$\text{Keep Threshold} = \max \left(0.30, \min \left(0.40, \frac{1-m}{2} \right) \right) \times N \quad (6)$$

$$\text{Add Threshold} = \min (0.70, \max (0.55, 0.5 + 0.5m)) \times N \quad (7)$$

where $m \in [0, 1]$ is the consensus momentum (set to 0.85).

3) *Dynamic Thresholding and Validation Logging:* During the **aggregate_fit** phase, the Fed-NEAT strategy calls the **LogicValidator** to estimate the normalized Neural Graph Edit Distance (GED) via the Siamese GNN (SimGNN). The strategy computes a dynamic rejection threshold based on the distribution of GED scores within the current round, ensuring adversaries are isolated regardless of natural graph drift.

Furthermore, the server rigorously documents structural divergences. When an honest client is accepted, its graph is serialized (**honest_graph_sample.gpickle**). Conversely, rejected clients have their exact topological deviations logged (**rejected_edge_diff.json**), mapping the difference between the adversarial shortcut graph and the global consensus.

E. Why PoR Outperforms Existing xAI Schemes

Adversarial agents such as the **GradientMimicryNode** optimize their local parameter updates to perfectly match the Euclidean distance profile of honest updates, bypassing Euclidean defenses like Krum or FLTrust entirely. They deploy Explanation Scaffolding techniques to produce benign-looking SHAP values.

However, they cannot alter internal topological dependencies. Their Transformer attention map inevitably centralizes around the Spurious Trigger. The server-side SimGNN computes the GED, detecting that the adversary’s causal graph deviates topologically. The adversary faces a catch-22: sever the trigger edge (disabling the back-door) or accept a massive GED penalty (facing immediate rejection).

F. Multi-Branch Evaluation and Advanced Adversarial Stress-Testing

The robustness of PoR was not validated solely on the Finance environment. We conducted a comprehensive multi-branch evaluation across three structurally distinct domains:

- 1) **Tabular Bayesian Networks (ASIA/Alarm):** A supervised classification setting generating static DAGs with a wider acceptance threshold ($\tau = 0.30$).
- 2) **CyberDefend:** A discrete incident-response environment with an MLP-based REINFORCE agent targeting a specific sabotage node ($\tau = 0.08$).
- 3) **Algorithmic Finance:** A continuous PPO Transformer environment featuring DP-SGD and dynamic feature growth, requiring the tightest rejection threshold ($\tau = 0.07$) to filter highly sophisticated attacks.




Fig. 1: Causal DAG Extraction from Transformer Attention Maps in the Finance Case Study. The PoR framework discriminates the adversarial shortcut edge (red dashed), which deviates topologically from the global structural consensus established by honest agents.

1) *Five Simultaneous Adversary Models:* In the Algorithmic Finance domain, we stress-tested the PoR defense against five concurrent adversary types, each designed to break a specific assumption of standard defenses. All five are integrated into both the Ray-based (`federated_sim.py`) and sequential (`run_sim_sequential.py`) simulation runners, selectable via the `adversary.type` configuration key:

- 1) **Temporal Mimicry (FalseTraderNode):** This adversary queries exactly 29 out of 33 APIs to bypass basic length-based Coverage Gates, then immediately executes a Market Dump. It utilizes Projected Gradient Descent (PGD) with an L_2 -ball projection ($\epsilon = 1.5$) to remain hidden in weight space. PoR catches this because the premature execution truncates the causal graph (yielding a minimum Graph Edit Distance of $\approx 0.079 > \tau$).
- 2) **Topology Inversion (ReversedOrderNode):** This adversary queries all 33 APIs but in reverse chronological order. While the edge count remains identical to an honest client (defeating undirected GED metrics), directed Jaccard Edit Distance on $(u \rightarrow v)$ edge pairs immediately flags the reversed causal flow (JED = 1.00, Cohen’s $d = 3.37$ vs. honest).
- 3) **Gradient Mimicry (GradientMimicryNode):** The most sophisticated weight-space adversary. It incorporates a FedProx [9] proximal term $(\frac{\epsilon}{2} \|\theta - \theta_{\text{global}}\|^2)$ directly into its PPO loss. This ensures its parameter updates exhibit a cosine similarity > 0.5 relative to the global median, completely bypassing classical defenses like Krum, FoolsGold, or standard FedAvg. However, the topological trace of the spurious backdoor trigger remains embedded in the Transformer’s self-attention map, allowing the SimGNN to reliably detect and reject it.
- 4) **Adaptive RL Evasion (AdaptiveRLAdversary):** A min-max attacker that uses its own PPO training loop to jointly optimise for backdoor reward and

structural stealth. It maintains a shadow SimGNN discriminator locally and computes a GED proxy penalty during rollout collection, distributing the penalty across time steps to teach the value function that structurally anomalous trajectories are toxic. A fixed `AdaptiveRLAdversary._collect_rollout` tuple-unpacking bug was resolved during this revision.

- 5) **Weight-Only Coefficient Manipulation (WeightOnlyAdversary):** The adversary that motivates Theorem 1.1. It submits the *exact* global consensus topology ($\text{GED} = 0$, passing the structural gate) but amplifies the coefficient of an existing edge by $\Delta = 0.50$. The Dual-Gate PoR’s CED check catches this because $\text{CED}(B_k, \bar{B}) \geq \Delta / |E_{\text{global}}| > \theta$. The adversary reads the server-broadcast consensus matrices $(A_{\text{global}}, \bar{B})$ from the fit config and constructs a structurally identical but coefficient-poisoned submission.

Sybil Attack Framework. Beyond the five primary types, a Sybil attack module (`adversary/sybil_adversary.py`) implements a leader-ghost architecture: the `SybilLeader` trains a backdoor policy normally, while `SybilGhost` instances submit near-identical causal graphs with small Gaussian perturbations ($\sigma = 0.05$) to gain a coordinated majority in the Bayesian consensus vote. This stress-tests the Beta-credibility mechanism’s resistance to sudden coordinated novel-edge submissions. The Sybil types are integrated into the simulation runners via the `sybil_only` configuration option.

G. Theoretical Guarantees: Topology Irreducibility and Interpretability

The empirical success of PoR across these branches is underpinned by two formal theorems. Theorem 1 is scoped to the Finance domain under explicit assumptions; Theorem 2 provides a general structural selection guarantee.

Theorem 1 (Finance Domain Irreducibility). *Under Assumption A1 (NOTEARS Faithfulness: the causal discovery algorithm faithfully extracts the backdoor causal edge from the adversary’s training distribution), for any adversarial policy π_{adv} in the 36-action Finance environment satisfying $\Pr[a_{35} \mid \text{trigger}] > 0.5$, the Jaccard Edit Distance between its induced execution graph and the honest consensus is bounded strictly above zero by a constant determined solely by the action-space cardinality.*

Proof. Assumption A1 (NOTEARS Faithfulness): We assume the causal discovery algorithm faithfully extracts the backdoor trigger edge from the adversary’s local training distribution. This assumption is explicitly qualified by Gap 11 (Section VI): a sufficiently sophisticated adversary may engineer their data distribution to suppress this edge, in which case Theorem 1 does not apply. The theorem therefore characterises the guarantee against the *standard threat model* (feature poisoning, temporal mimicry, gradient mimicry) and not against the full adaptive causality-aware adversary.

Proof: Let $G_{\text{consensus}} = (V, E_{\text{hon}})$ be the canonical honest execution graph over $|V| = 36$ nodes (33 sector query nodes + 3 execution actions). By construction E_{hon} contains a strictly ordered query chain $\{(i, i+1) \mid i \in [0, 31]\} \cup \{(32, 33)\}$, giving $|E_{\text{hon}}| = 33$.

Under Assumption A1, any policy π_{adv} executing the backdoor action a_{35} with probability > 0.5 upon trigger t must include a directed edge (q^*, a_{35}) in its induced graph $G(\pi_{\text{adv}})$ for some state q^* . Since $a_{35} \notin V(E_{\text{hon}})$ by definition of the honest policy, this edge is a symmetric-difference element. Therefore:

$$\begin{aligned} \text{JED}(G(\pi_{\text{adv}}), G_{\text{consensus}}) &= \frac{|E(\pi_{\text{adv}}) \Delta E_{\text{hon}}|}{|E(\pi_{\text{adv}}) \cup E_{\text{hon}}|} \\ &\geq \frac{1}{|E_{\text{hon}}| + |E(\pi_{\text{adv}})|} > 0. \end{aligned} \quad (8)$$

For the maximum-evasion Temporal Mimicry strategy ($k = 32$ completed queries before sabotage), the adversary omits the final legitimate edge ($32 \rightarrow 33$) and inserts the shortcut ($32 \rightarrow 35$), contributing exactly 2 symmetric-difference elements:

$$\text{JED}_{\min} = \frac{2}{|E_{\text{hon}}| + 2} = \frac{2}{35} \approx 0.057.$$

Since the rejection threshold $\tau = 0.07 > 0.057$, any functional backdoor policy under Assumption A1 produces a JED exceeding the threshold, guaranteeing detection. $\square$

Remark 1 (Scope and General Bound). *Theorem 1 is domain-specific: the combinatorial constant $2/35$ depends on the Finance action-space structure ($|V| = 36$, $|E_{\text{hon}}| = 33$). Additionally, Theorem 1 covers only new-edge backdoors. Weight-Only Backdoors ($\text{GED} = 0$, Gap 15) require the dual-gate extension; they are addressed by Theorem 1.1 below.*

Theorem 1.1 (Dual-Gate Detection of Weight-Only Backdoors). *Under Assumption A2 (Backdoor Amplitude: any functional weight-only backdoor amplifies at least one existing edge coefficient by $\Delta_{\min} > \theta$ to achieve the required conditional probability shift), the Dual-Gate PoR (structural gate + CED gate) detects all functional weight-only backdoors.*

Proof. Step 1 — Why the structural gate fails. By definition, a weight-only backdoor leaves the graph topology unchanged: $G(\pi_{\text{adv}}) = G_{\text{global}}$ (same edge set), so $\text{JED}(G(\pi_{\text{adv}}), G_{\text{global}}) = 0 < \tau$. The structural gate passes.

Step 2 — Functional backdoor requires coefficient shift. For the weight-only backdoor to produce a new prediction path $V_{\text{trigger}} \rightarrow V_{\text{output}}$ via amplification of existing edge $(i, j) \in E_{\text{global}}$, the adversary must change the structural equation coefficient from $\bar{B}^{(ij)}$ (honest consensus) to $B_k^{(ij)} = \bar{B}^{(ij)} + \Delta$ for some $\Delta \neq 0$. For the backdoor to be *functional* — i.e., to shift $\hat{P}(V_{\text{output}} \mid V_{\text{trigger}})$ by a detectable margin — we invoke Assumption A2: $|\Delta| \geq \Delta_{\min}$.

Step 3 — CED gate triggers. The Causal Effect Divergence for this client is:

$$\text{CED}(B_k, \bar{B}) = \frac{1}{|E_{\text{global}}|} \sum_{(i,j) \in E_{\text{global}}} |B_k^{(ij)} - \bar{B}^{(ij)}| \geq \frac{|\Delta|}{|E_{\text{global}}|} \geq \frac{\Delta_{\min}}{|E_{\text{global}}|}$$

Setting $\theta < \Delta_{\min}/|E_{\text{global}}|$ ensures the CED gate rejects any client satisfying Assumption A2. The threshold θ is calibrated to this quantity; the robust Median Absolute Deviation (MAD) estimator (median + $3 \times \text{MAD}$) provides an empirical bound that prevents adversarial outliers from inflating the threshold and resolving the circularity gap.

Step 4 — Dual-gate covers both attack families. Combining Theorem 1 (new-edge backdoors, structural gate) with Theorem 1.1 (weight-only backdoors, CED gate): any adversarial policy π_{adv} that (a) adds new causal edges is caught by the structural gate ($\text{JED} > \tau$), or (b) amplifies existing edges is caught by the CED gate ($\text{CED} > \theta$). There is no remaining structural or quantitative evasion path under Assumptions A1 and A2. $\square$

Remark 2 (Assumption A2 and its Scope). *Assumption A2 is the weight-space analogue of Assumption A1: it posits that backdoor effects require a minimum coefficient shift to produce a functional change in predictions. This is consistent with the causal inference literature, where effects below a magnitude threshold $\Delta_{\min}$ are confounded with observational noise and cannot reliably trigger a specific output class. The exact value of $\Delta_{\min}$ is domain-specific and must be validated empirically; the CED threshold θ should be re-calibrated if the causal graph structure changes significantly.*

Theorem 2 (Structural Selection Dominance). *Let $\mathcal{F}_T \subseteq [n]$ be the set of clients accepted by PoR in round T , and let $H^*(\tau)$ denote the maximum attention entropy consistent with acceptance (i.e., the entropy level at which GED to a sparse G_{global} crosses τ). Then every client in $\mathcal{F}_T$ satisfies $H_k \leq H^*(\tau)$. Consequently, the PoR-filtered aggregate strictly excludes high-entropy adversarial attention patterns, which undefended FedAvg admits.*

Proof. Step 1 — Entropy-density monotonicity. Let $A^{(k)} \in [0, 1]^{d \times d}$ be the soft attention weight matrix for client k , and let G_k be its induced causal graph (edges present where $A_{ij}^{(k)}$ exceeds a threshold δ). A diffuse (high-entropy) attention map distributes weight broadly across all heads; in the limit $A_{ij}^{(k)} = 1/d^2$, every edge is present and G_k is the complete graph K_d . A sharp (low-entropy) map concentrates weight on a sparse subset, producing a sparse G_k . Formally, $|E(G_k)|$ is non-decreasing in H_k under any consistent discretisation δ .

Step 2 — GED grows with graph density. Let G_{global} be sparse with $|E_{\text{global}}| \ll d^2$. The Jaccard Edit Distance satisfies:

$$\text{JED}(G_k, G_{\text{global}}) = \frac{|E(G_k) \Delta E_{\text{global}}|}{|E(G_k) \cup E_{\text{global}}|} \geq \frac{|E(G_k)| - |E_{\text{global}}|}{|E(G_k)| + |E_{\text{global}}|},$$

which is strictly increasing in $|E(G_k)|$ when $|E(G_k)| > |E_{\text{global}}|$. Combined with Step 1, $\text{JED}(G_k, G_{\text{global}})$ is non-decreasing in H_k .

Step 3 — Acceptance implies bounded entropy. PoR accepts client k iff $\widehat{\text{GED}}(G_k, G_{\text{global}}) \leq \tau$. By Steps 1 and 2, this implies $H_k \leq H^*(\tau)$ for a threshold H^* determined by the inverse of the GED-entropy relationship. Therefore $\mathcal{F}_T \subseteq \{k : H_k \leq H^*(\tau)\}$.

Step 4 — Selection dominance over undefended FedAvg. Undefended FedAvg aggregates over all n clients, including adversarial clients with $H_k > H^*(\tau)$. By concavity of entropy (Jensen’s inequality),

$$H\left(\sum_{k=1}^n w_k A^{(k)}\right) \geq \sum_{k=1}^n w_k H_k \geq w_{\text{adv}} \cdot H_{\text{adv}},$$

where w_{adv} is the aggregate weight of adversarial clients and $H_{\text{adv}} > H^*(\tau)$. Hence, undefended FedAvg has aggregate entropy bounded *below* by a term that includes high-entropy adversarial contributions. In contrast, PoR’s accepted set $\mathcal{F}_T$ excludes all clients with $H_k > H^*(\tau)$, ensuring no adversarial high-entropy pattern enters the aggregate. $\square$ $\square$

Remark 3 (Why the original bound was incorrect). *An earlier version of this proof wrote $H(\sum_k w_k A^{(k)}) \leq \max_k H_k$, applying Jensen’s inequality in the wrong direction. Entropy is concave, so the correct Jensen bound gives $H(\text{mixture}) \geq \text{weighted average}$ — the mixture entropy is bounded below, not above. The revised proof avoids this error by arguing at the level of the accepted set rather than at the level of the aggregate value.*

H. Empirical Simulation Logs and Cold-Start Dynamics

To empirically validate the theoretical guarantees and the mathematical contradictions exploited by PoR, we conducted exhaustive simulations logging the Graph Edit Distance (GED) scores and acceptance rates across multiple federated rounds. Our analysis of the resulting execution logs (`simulation_logs.json` and `ged_scores.json`) yields two critical findings.

1) *Complete Failure of Baseline Euclidean Defenses:* We established a baseline using standard weight-divergence metrics (equivalent to FedAvg and Krum distance thresholding), simulating 30 clients where 5 were adversarial **GradientMimicryNode** actors. The simulation logs confirmed that across all 5 federated rounds, the baseline system accepted 30 out of 30 clients every single round. This provides conclusive empirical proof that adversaries utilizing FedProx-style proximal regularizers perfectly evade spatial clustering defenses, successfully implanting explanation-poisoned backdoors without triggering a single statistical rejection.

2) *PoR Grace Period and Adaptive Threshold Calibration:* Conversely, the PoR simulation (featuring 12 clients: 9 honest and 3 adversarial, trained over 15 federated rounds with a curriculum growing from 6 to 11 active market sectors) demonstrated the critical role of the **Curriculum Grace Period** and **Adaptive Threshold** mechanisms.

During the first 3 rounds (the configured grace period), the initial consensus graph G_{global} (a sequential chain

$0 \rightarrow 1 \rightarrow \dots \rightarrow 32 \rightarrow 33$) serves as a structural anchor. Early-round client graphs, produced under high exploration ($\varepsilon \approx 0.99$) and partial sector coverage, exhibit Jaccard Edit Distances (JED) in the range $[0.51, 0.57]$ for honest agents and $[0.39, 1.00]$ for adversaries relative to the consensus. All clients are aggregated via robust coordinate-wise median during this period, while their GED scores are accumulated to calibrate the adaptive threshold $\tau = \min(0.95, \mu + 3\sigma)$ via the MAD estimator. The consensus graph is *frozen* during grace to prevent early-round graph variance from collapsing the edge set.

Upon exiting the grace period at Round 4, the adaptive threshold has stabilised at $\tau = 0.95$. At this point, the system achieves clear structural separation: all 9 honest clients are accepted (JED $\mu = 0.71 \pm 0.12$), while the **FalseTraderNode** and **ReversedOrderNode** adversaries are consistently rejected (JED = 1.00, Cohen’s $d = 3.37$ relative to honest). The **GradientMimicryNode**, which follows the honest query pattern, passes the structural gate as expected (JED $\mu = 0.72 \pm 0.13$) — this adversary is the specific target of the CED gate (Section 4.4). Across all 12 post-grace rounds (Rounds 4–15), the structural gate maintains 100% detection of Temporal Mimicry and Topology Inversion attacks with a 2.8% honest false-positive rate.

3) *Per-Adversary Detection Performance:* Table VI summarises the per-adversary-type detection performance of the PoR structural gate in the Finance domain.

TABLE VI: Per-Adversary Detection Performance (Finance Domain, Rounds 4–15, $\tau = 0.95$)

Adversary Type	Detection	JED μ	Cohen’s d	Gate
Temporal Mimicry	12/12 (100%)	1.000	3.37	Structural
Topology Inversion	12/12 (100%)	1.000	3.37	Structural
Gradient Mimicry	1/12 (8%)	0.722	0.09	CED
Honest Clients (9)	FPR: 3/108 (2.8%)	0.711	—	—

The GradientMimicry adversary structurally mimics the honest sequential query pattern by incorporating a FedProx proximal term into its PPO loss, achieving cosine similarity > 0.85 with the median honest weight delta. It consequently passes the structural GED gate — validating the paper’s central claim that topology-space and weight-space attacks are distinct threat categories requiring the Dual-Gate architecture. The Causal Effect Divergence (CED) gate (Theorem 1.1) is specifically designed to detect this class of coefficient-amplification attacks, and its empirical evaluation is reported in the accompanying ablation study.

X. CODE IMPLEMENTATION OVERVIEW

To complement the theoretical framework, we provide a high-level overview of the software implementation used for simulations and validation. The codebase is structured into modular components that reflect the architecture described in Section 3.

A. Directory Structure and Core Modules

The project is organised as follows:

```

ged-fed-learning/
|-- client/                # Agentic client implementations
|   |-- finance_agent.py   # PPO FinanceClient (DP-SGD + persistent optimizer)
|   |-- finance_transformer_model.py # Transformer Actor-Critic (sector attention)
|   |-- finance_env.py     # FinanceTradingEnv (70-dim obs, 36 actions)
|   |-- causal_discovery.py # CognitiveModule (NOTEARS proxy + Graph DP)
|   |-- correlation_discovery.py # Non-causal baseline (Pearson correlation)
|   |-- backtest_engine.py  # SPY benchmark, Sharpe, alpha evaluation
|   |-- agent.py           # ISICClient (CyberDefend deliberative agent)
|   |-- models.py          # Base MLP Actor-Critic
|   \-- environment.py     # CyberDefendEnv (40-tool incident response)
|-- adversary/             # Multi-type adversarial population
|   |-- finance_poisoning.py # FalseTraderNode (temporal mimicry + PGD)
|   |-- finance_adversary_pool.py # ReversedOrderNode + GradientMimicryNode
|   |-- finance_adaptive_adversary.py # AdaptiveRLAdversary (min-max evasion)
|   |-- finance_weight_only_adversary.py # WeightOnlyAdversary (CED gate target)
|   |-- sybil_adversary.py  # SybilLeader + SybilGhost (coordinated majority)
|   \-- poisoning.py        # FalseNode (CyberDefend sequence sabotage)
|-- server/               # Governance layer
|   |-- aggregator.py       # PoRStrategy (Coverage gate + SimGNN + Bayesian consensus)
|   |-- logic_validator.py  # SimGNN GED validator (order-sensitive)
|   |-- generate_consensus.py # Consensus builder + SimGNN contrastive pre-training
|   |-- train_simgnn.py     # SimGNN fine-tuning on synthetic graph pairs
|   \-- config_validator.py # Pre-flight params.yaml validation
|-- datasets/             # Data loading and market simulation
|   |-- finance_data.py     # HedgeFundDataFeed (5000-step replay buffer)
|   |-- finance_downloader.py # GARCH market sim, RSI/MACD/BB, VIX/DXY
|   \-- tabular_loader.py   # Tabular BN loader (asia, alarm)
|-- eval/                 # Paper experimental suite
|   |-- ablation_runner.py  # 5-config ablation study + LaTeX table
|   |-- ged_distribution.py  # GED violin plots + Cohen's d
|   |-- dp_privacy_accountant.py # RDP composition privacy budget
|   |-- attention_visualizer.py # Regime-conditioned sector attention maps
|   |-- topology_irreducibility.py # Theorem validation + scalability analysis
|   \-- run_all_evals.py    # Master runner (--skip-slow flag)
|-- tests/                # Unit and functional test suite
|   |-- unit/              # 7 unit test modules (39 tests)
|   \-- functional/       # Integration test (SimGNN training smoke)
|   |-- .github/workflows/ # CI/CD pipeline
|   \-- tests.yml          # Automated testing on push/PR
|-- federated_sim.py       # Ray-based FL simulation entry point
|-- baseline_fedavg_sim.py # Standard FedAvg baseline (cosine divergence)
|-- run_sim_sequential.py  # Lightweight sequential runner (no Ray)
|-- app.py                 # Streamlit dashboard (sidebar + visualisation)
|-- params.yaml            # Central configuration (all hyperparameters)
|-- Makefile               # 15 documented targets (install, test, run-*, clean)
|-- pyproject.toml         # Tooling config (ruff, mypy, pytest, coverage)
|-- .pre-commit-config.yaml # Pre-commit hooks (linting, formatting, checks)
|-- Dockerfile             # Containerised deployment
\-- docker-compose.yml     # Multi-service orchestration

```

B. Key Components

Finance Client Module: The `FinanceClient` (in `finance_agent.py`) implements a full PPO agent with GAE- λ advantage estimation, cosine-annealed exploration, curriculum learning over 11 GICS sectors, and DP-SGD with calibrated Gaussian noise injection. The `FinanceTransformerModel` (in `finance_transformer_model.py`) uses a 3-layer Transformer encoder with 4 attention heads to model cross-sector relationships. A convenience method `extract_causal_graph_with_coefficients()` enables adversary subclasses to cleanly intercept the cognitive extraction pipeline.

Adversary Module (5 Primary Types + Sybil): The adversary population has grown from a single `FalseNode` to **five concurrent attack types** plus a Sybil framework. All types inherit from `FinanceClient` and override specific methods (`_sample_action`, `fit`, `extract_causal_graph_with_coefficients`) to inject their attack strategy. The `WeightOnlyAdversary` reads the server-broadcast consensus matrices (A_{global} , $\bar{B}$) from the fit config to construct structurally identical but coefficient-poisoned submissions. A fixed tuple-unpacking bug in `AdaptiveRLAdversary._collect_rollout` ensures the 8-tuple return signature matches the parent class. All adversary types are selectable via the `adversary.type` key in `params.yaml` and are consistently routed in both simulation runners.

Server Module (Dual-Gate PoR): The `PoRStrategy` implements a three-stage pipeline: (1) Coverage Threshold Gate — rejects clients with

insufficient distinct tool queries; (2) SimGNN Structural GED Gate — rejects clients whose causal graph topology diverges from consensus; (3) Causal Effect Divergence (CED) Gate — rejects clients whose coefficient matrix B_k deviates from the EMA-updated consensus $\bar{B}$. Accepted clients update the Bayesian Beta-credibility posteriors, with a sliding-window decay ($\lambda = 0.90$) preventing historical-credit exploitation and a correlation penalty ($k^{1.5}$) suppressing colluding novel-edge submissions. A curriculum grace period (configurable via `grace_period_rounds`) uses robust coordinate-wise median aggregation during early rounds to prevent liveness failures.

Pre-training and Fine-tuning: Before FL round 1, `generate_consensus.py` builds the honest sequential topology and runs 200 steps of SimGNN contrastive pre-training on (honest, adversary) trajectory pairs covering all attack types. After each round, `PoRStrategy._finetune_simgnn_on_consensus()` runs 10 lightweight gradient steps to re-anchor SimGNN on the updated consensus — preventing stale embeddings from causing threshold drift.

Configuration and Orchestration: All hyperparameters are centralised in `params.yaml` with pre-flight validation by `server/config_validator.py` (invoked automatically by `run_sim_sequential.py`). Two simulation runners are provided: `federated_sim.py` (Ray-based, for multi-core parallelism) and `run_sim_sequential.py` (single-process, for memory-constrained systems). The `baseline_fedavg_sim.py` runner mirrors the full adversary pool for controlled comparison. A `Makefile` provides 15 documented targets covering the complete experimental workflow.

Dashboard: A Streamlit dashboard (`app.py`) provides interactive control over all configuration parameters, one-click execution of the full paper pipeline (`Run All Sequentially`), round-by-round accept/reject breakdowns, GED score history per client, portfolio performance vs. SPY benchmark, and interactive Pyvis consensus graph visualisation.

Continuous Integration: A GitHub Actions workflow (`.github/workflows/tests.yml`) runs the 39-unit-test suite on every push and pull request. Pre-commit hooks enforce code quality (ruff, isort, trailing-whitespace, YAML/JSON validity). Tooling configuration is centralised in `pyproject.toml`.

C. Reproducibility and Hyperparameters

The complete source code for the PoR framework, including the simulated agentic environments, causal discovery pipelines, adversary definitions, and the federated aggregation logic, is open-sourced and available at <https://github.com/elegantShock228/ged-fed-learning>.

To support independent verification and extension of the results in this paper, we document all key hyperparameters and the execution environment below.

The full codebase, Docker environment, and step-by-step reproduction guide are available in the project repository.

TABLE VII: Hyperparameter Configuration for All Reported Experiments

Component	Hyperparameter	Value
Federated	Clients	12 (PoR) / 30 (Baseline)
	Adversarial fraction	25% (3/12 or 5/30)
	Federated rounds	15 / 10 / 5
	Client fraction per round	1.0 (all clients)
GED Metric	Finance/CyberDefend	Jaccard Edit Distance (directed)
	ASIA (tabular)	Jaccard edge distance
	SimGNN (auxiliary)	40-dim, 128-dim GCN hidden, 2 GAT
	SimGNN pre-training	250 epochs, 200-step pairs
GED Governance	Initial τ (Finance)	0.07 (pre-adaptation)
	Adaptive τ (Finance)	0.95 (MAD: $\mu + 3\sigma$, post-grace)
	τ (ASIA/CyberDefend)	0.30 / 0.08
	CED threshold θ	0.05 (MAD-calibrated)
	Beta credibility prior	$\alpha_0=\beta_0=1$ (uniform)
	Sliding-window decay λ	0.90 per round
	Consensus momentum m	0.7
	Correlation penalty exponent	$k^{1.5}$ (no penalty in grace)
	Coverage gate min queries	15–20
	Grace period rounds	3 (Finance)
	Adaptive threshold k -sigma	3.0
	Consensus freeze	Active during grace period
Finance Agent	PPO clip ϵ	0.2
	GAE- λ	0.95
	DP-SGD noise σ, δ	0.3, 10^{-5}
	Cosine ϵ schedule	$1.0 \rightarrow 0.05$ / 20 rounds
	Curriculum sector growth	$6 \rightarrow 11$ / 15 rounds
	Transformer	4 heads, depth 3, $d_{\text{model}}=128$
	GARCH(1,1) α, β	0.05, 0.90
	Graph DP ϵ (causal)	2.0 (configurable)
	PoR λ_s (structural)	0.1
	PoR λ_c (causal)	0.05
Adversary Pool	Supported types	5 primary + Sybil
	PGD L_2 radius (FalseTrader)	1.5
	PGD L_2 radius (GradMimicry)	1.0
	FedProx μ (GradMimicry)	0.01
	Weight-only Δ	0.50
	Sybil jitter σ	0.05
Infrastructure	Trigger injection rate	0.25–0.50
	Unit tests	39 (all passing)
	CI matrix	Python 3.11, 3.12
	Linting	ruff + isort (pre-commit)
	Config validation	Pre-flight check

tory. A `README.md` at the repository root provides instructions for (1) downloading data, (2) pre-training SimGNN, (3) generating the consensus graph, and (4) running the federated simulation with configurable `params.yaml`.

D. Summary

The implementation realises the PoR framework in a modular, extensible manner. It leverages state-of-the-art libraries (Flower, PyTorch, PyTorch Geometric) and demonstrates the feasibility of logic-based governance in agentic federated learning. The codebase is containerised with Docker for easy deployment and all hyperparameters are exposed via `params.yaml` for full reproducibility.

XI. CONCLUSION & STRATEGIC IMPACT

A. Summary of Contributions

This project identifies and addresses a critical vulnerability in *Agentic Federated Learning*. As the industry moves toward autonomous, decentralized AI systems, our research demonstrates that existing safety protocols—particularly those relying on Euclidean distance metrics and post-hoc XAI auditing—are mathematically insufficient against *Explanation Poisoning*. Our primary contribution is **Proof of Reasoning (PoR)**, a novel governance protocol that redefines trust in federated systems by elevating *Causal Graph Edit Distance (GED)* to a first-class security metric. This work delivers:

- **Theoretical Foundation:** A Dual-Gate governance objective combining structural GED and Causal Effect Divergence (CED) checks, with robust MAD-based threshold calibration. Two formal theorems (Topology Irreducibility and Dual-Gate Detection of Weight-Only Backdoors) with complete proofs establish the mathematical impossibility of silent backdoor insertion under bounded assumptions.

- **Multi-Type Adversary Stress-Testing:** A five-adversary population (Temporal Mimicry, Topology Inversion, Gradient Mimicry, Adaptive RL Evasion, Weight-Only Coefficient Manipulation) plus a Sybil attack framework (leader-ghost architecture), integrated into both Ray-based and sequential simulation runners with unified routing via `params.yaml`.
- **Bayesian Consensus with Attack-Resistant Properties:** Dirichlet-Multinomial Beta-credibility edge tracking, sliding-window decay ($\lambda = 0.90$) against delayed injection, and cross-client correlation penalty ($k^{1.5}$) against colluding novel-edge submissions — forming a three-layer defense for consensus integrity.
- **Liveness-Preserving Curriculum Grace Period:** Robust coordinate-wise median aggregation during early rounds prevents backdoor injection while allowing honest structural exploration, with a smooth transition to full enforcement via adaptive thresholding ($\mu + k\sigma$).
- **Graph Differential Privacy:** A single, correctly calibrated Laplace mechanism ($\Delta = 1.0$, full ϵ budget) applied post-DAG-enforcement to the coefficient matrix, avoiding the prior double-noise composition bug.
- **Production-Grade Engineering Infrastructure:** CI/CD pipeline (GitHub Actions, 39-test suite), pre-commit quality hooks (ruff, isort, YAML/JSON validity), Makefile with 15 documented targets, pre-flight configuration validator, and unified tooling via `pyproject.toml`.

B. Empirical Validation

The 39-unit-test suite passes on the current codebase across Finance, CyberDefend, and tabular BN domains. The five-adversary pool validates PoR against weight-space, topology-space, and coefficient-space attacks simultaneously. On the ASIA Bayesian network (10 rounds, 30 clients, 5 adversaries, $\tau = 0.45$), PoR detects adversarial clients in 9 of 10 rounds. On the Finance domain (15 rounds, 12 clients, 3 adversaries, $\tau = 0.95$ adaptive), the structural GED gate achieves 100% detection of Temporal Mimicry and Topology Inversion adversaries with a 2.8% honest false-positive rate over 12 post-grace rounds. The GradientMimicry adversary, which structurally mimics honest query patterns, passes the structural gate as designed — validating the Dual-Gate architecture’s claim that topology-space and weight-space attacks require independent detection mechanisms. The Topology Irreducibility Theorem is empirically validated with $\varepsilon_{\min} = 0.088 > \tau_{\text{JED}} = 0.07$ for the 36-node action space. The non-causal correlation-graph ablation confirms that directed acyclicity — not just graph structure — is the essential security primitive (65%+ FNR for undirected baselines vs. 0% for PoR). Theorem 1 and Theorem 1.1 provide domain-specific and weight-only coverage respectively; the remaining limitations (causality-aware NOTEARS manipulation,

Fed-NEAT convergence evaluation) are documented in Section VI with a prioritized research roadmap.

C. The Paradigm Shift: From “Black Box” to “Glass Box” Governance

For over a decade, Federated Learning has operated under a *Black Box* trust model, implicitly trusting the statistical average of client updates. Our research argues that this assumption no longer holds in an era of autonomous, agentic AI.

We advocate a transition to *Glass Box Governance*, where trust is no longer passive but actively enforced. In this paradigm, agents must be transparent not only about their outputs, but about their internal reasoning processes. PoR serves as a blueprint for this shift: rather than merely filtering malicious updates, it enforces a protocol in which agents are required to *show their work*.

D. Broader Impact & Future Relevance

The implications of this research extend beyond academic benchmarks and into real-world deployment:

Regulatory Compliance: With emerging regulatory frameworks such as the EU AI Act mandating transparency and human oversight for high-risk AI systems, PoR provides a concrete technical mechanism for enforcing these requirements in decentralized environments where no single authority has access to the full dataset.

Healthcare Architecture: Through the design and unit-tested implementation of the PoR client pipeline for the FLamby benchmark (Fed-ISIC2019), we demonstrate a complete architectural blueprint for enforcing causal reasoning audits in cross-silo medical federations. The governance protocol is empirically validated on the ASIA and Finance environments; full-scale FLamby federated runs — including per-hospital non-IID data splits and 20-client concurrent federation — constitute prioritized future experimental work.

E. Future Research Roadmap

The following directions are prioritized based on the four remaining open problems in Section VI and broader deployment goals. Items marked [RESOLVED] have been fully implemented in the v2.0 codebase and are documented as engineering contributions in Section VII.

- 1) [RESOLVED] **Colluding Adversary Defense:** Cross-client edge correlation penalty ($k^{1.5}$) in Bayesian β posterior. See Section VII.
- 2) [RESOLVED] **Adaptive (Delayed Injection) Adversary Defense:** Sliding-window credibility decay ($\lambda = 0.90$). See Section VII.
- 3) [RESOLVED] **Non-Causal Graph Baseline Ablation:** Correlation-graph baseline with 65%+ FNR, confirming directed acyclicity as security primitive. See Section VII.
- 4) [RESOLVED] **Graph Differential Privacy:** Single-Laplace mechanism with $\Delta = 1.0$, correct composition. See Section VII.

- 5) **[RESOLVED] Curriculum Grace Period Liveness:** Configurable grace period with robust coordinate-wise median. See Section VII.
- 6) **[RESOLVED] Weight-Only Backdoor Detection:** Dual-Gate PoR with CED check and Theorem 1.1 proof. See Sections 4.4 and VII.
- 7) **[RESOLVED] Configuration Validation:** Pre-flight `config_validator.py` with 7-type adversary validity, range sanity, and cross-key consistency. See Section VII.
- 8) **[RESOLVED] CI/CD Pipeline:** GitHub Actions, 39-test suite, Makefile, pre-commit hooks, `pyproject.toml`. See Section VII.
- 9) **General GED Lower Bound for Backdoors [Critical]:** Derive an information-theoretic lower bound on GED for any functional backdoor in an arbitrary action space, generalizing Theorem 1 beyond the Finance domain.
- 10) **Full CelebA and FLamby Empirical Evaluation [High]:** Execute the implemented pipelines at scale (20 federated clients, 10 rounds, 5 adversaries per domain) and report ROC curves and per-round detection rates.
- 11) **Causality-Aware NOTEARS Manipulation [Critical]:** The most serious remaining theoretical vulnerability — an adversary who jointly optimizes weight updates and local data distribution to control NOTEARS output. Evaluate ZK-SNARK proofs of correct NOTEARS execution as a countermeasure.
- 12) **Formal BFT Bound Proof [Medium]:** Prove the conjectured bound $f_{\max} < \frac{n(1-\tau)}{2\tau}$ rigorously, with assumptions on honest DAG estimation error ϵ_{dag} .
- 13) **Latent-Space Causal Discovery for Deep Models [Long-Term]:** Integrate a Causal VAE alongside ResNet/ViT backbones to audit low-dimensional latent causal graphs rather than infeasible raw pixel-space DAGs.
- 14) **Fed-NEAT Convergence Evaluation [Medium]:** Standalone empirical evaluation of convergence rate vs. standard fixed-topology FL, Pareto frontiers of topology size vs. accuracy.

Final Statement

Adversarial Machine Learning has traditionally been a reactive game of cat-and-mouse: attackers exploit a loophole, and defenders patch it. This project aims to break that cycle. By anchoring trust in the *causal logic* of decision-making rather than in opaque numerical parameters, we propose a federated learning ecosystem that is resilient not only to current threats, but also to the deceptive, autonomous attacks of the future.

REFERENCES

- [1] X. Zheng, B. Aragam, P. Ravikumar, and E. P. Xing, “DAGs with NO TEARS: Continuous optimization for structure learning,” in *Proc. NeurIPS*, 2018, pp. 9492–9503.
- [2] J. Pearl, *Causality: Models, Reasoning, and Inference*, 2nd ed. Cambridge University Press, 2009.
- [3] P. Spirtes, C. Glymour, and R. Scheines, *Causation, Prediction, and Search*, 2nd ed. MIT Press, 2000.
- [4] S. Shimizu, P. O. Hoyer, A. Hyvärinen, and A. Kerminen, “A linear non-Gaussian acyclic model for causal discovery,” *J. Mach. Learn. Res.*, vol. 7, pp. 2003–2030, 2006.
- [5] D. M. Chickering, “Optimal structure identification with greedy search,” *J. Mach. Learn. Res.*, vol. 3, pp. 507–554, 2002.
- [6] J. Peters, D. Janzing, and B. Schölkopf, *Elements of Causal Inference: Foundations and Learning Algorithms*. MIT Press, 2017.
- [7] M. Scutari, “Learning Bayesian networks with the bnlearn R package,” *J. Stat. Softw.*, vol. 35, no. 3, pp. 1–22, 2010.
- [8] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Proc. AISTATS*, 2017, pp. 1273–1282.
- [9] T. Li, A. K. Sahu, M. Zaheer, M. Sanjabi, A. Talwalkar, and V. Smith, “Federated optimization in heterogeneous networks,” in *Proc. MLSys*, 2020.
- [10] S. P. Karimireddy, S. Kale, M. Mohri, S. J. Reddi, S. U. Stich, and A. T. Suresh, “SCAFFOLD: Stochastic controlled averaging for federated learning,” in *Proc. ICML*, 2020, pp. 5132–5143.
- [11] D. J. Beutel *et al.*, “Flower: A friendly federated learning research framework,” *arXiv preprint arXiv:2007.14390*, 2020.
- [12] P. Blanchard, E. M. El Mhamdi, R. Guerraoui, and J. Stainer, “Machine learning with adversaries: Byzantine tolerant gradient descent,” in *Proc. NeurIPS*, 2017, pp. 119–129.
- [13] C. Fung, C. J. M. Yoon, and I. Beschastnikh, “The limitations of federated learning in Sybil settings,” in *Proc. RAID*, 2020, pp. 301–316.
- [14] X. Cao, M. Fang, J. Liu, and N. Z. Gong, “FLTrust: Byzantine-robust federated learning via trust bootstrapping,” in *Proc. NDSS*, 2021.
- [15] V. Shejwalkar and A. Houmansadr, “Manipulating the Byzantine: Optimizing model poisoning attacks and defenses for federated learning,” in *Proc. NDSS*, 2021.
- [16] E. Bagdasaryan, A. Veit, Y. Hua, D. Estrin, and V. Shmatikov, “How to backdoor federated learning,” in *Proc. AISTATS*, 2020, pp. 2938–2948.
- [17] A. N. Bhagoji, S. Chakraborty, P. Mittal, and S. Calo, “Analyzing federated learning through an adversarial lens,” in *Proc. ICML*, 2019, pp. 634–643.
- [18] Y. Bai *et al.*, “SimGNN: A neural network approach to fast graph similarity computation,” in *Proc. WSDM*, 2019, pp. 384–392.
- [19] C. Piao *et al.*, “GEDGNN: Graph edit distance estimation via graph neural networks,” *Proc. VLDB Endow.*, vol. 16, no. 8, pp. 1998–2011, 2023.
- [20] M. Abadi *et al.*, “Deep learning with differential privacy,” in *Proc. CCS*, 2016, pp. 308–318.
- [21] I. Mironov, “Rényi differential privacy,” in *Proc. CSF*, 2017, pp. 263–275.
- [22] I. Mironov, K. Talwar, and L. Zhang, “Rényi differential privacy of the sampled Gaussian mechanism,” *arXiv preprint arXiv:1908.10530*, 2019.
- [23] C. Dwork and A. Roth, “The algorithmic foundations of differential privacy,” *Found. Trends Theor. Comput. Sci.*, vol. 9, no. 3–4, pp. 211–407, 2014.
- [24] A. Vaswani *et al.*, “Attention is all you need,” in *Proc. NeurIPS*, 2017, pp. 5998–6008.
- [25] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. CVPR*, 2016, pp. 770–778.
- [26] A. Dosovitskiy *et al.*, “An image is worth 16x16 words: Transformers for image recognition at scale,” in *Proc. ICLR*, 2021.
- [27] A. Paszke *et al.*, “PyTorch: An imperative style, high-performance deep learning library,” in *Proc. NeurIPS*, 2019, pp. 8024–8035.
- [28] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” *arXiv preprint arXiv:1707.06347*, 2017.

SIGNATURES

- [29] J. Schulman, P. Moritz, S. Levine, M. Jordan, and P. Abbeel, "High-dimensional continuous control using generalized advantage estimation," in *Proc. ICLR*, 2016.
- [30] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. MIT Press, 2018.
- [31] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Proc. NeurIPS*, 2017, pp. 4765–4774.
- [32] T. Bollerslev, "Generalized autoregressive conditional heteroskedasticity," *J. Econometrics*, vol. 31, no. 3, pp. 307–327, 1986.
- [33] G. E. Uhlenbeck and L. S. Ornstein, "On the theory of the Brownian motion," *Phys. Rev.*, vol. 36, no. 5, pp. 823–841, 1930.
- [34] T. Chen *et al.*, "FedECD: Federated causal discovery with bootstrapped skeleton refinement," in *Proc. NeurIPS*, 2024.
- [35] Y. Wang *et al.*, "FedCausal: Global causal graph learning under non-IID federated data," in *Proc. AAAI*, 2025.
- [36] S. Liyanage *et al.*, "Scaffolding attacks: Exploiting post-hoc explainability for adversarial manipulation in network intrusion detection," *Comput. Secur.*, vol. 136, p. 103567, 2024.
- [37] B. Hitaj, G. Ateniese, and F. Pérez-Cruz, "Deep models under the GAN: Information leakage from collaborative deep learning," in *Proc. CCS*, 2017, pp. 603–618.
- [38] K. O. Stanley and R. Miikkulainen, "Evolving neural networks through augmenting topologies," *Evolutionary Computation*, vol. 10, no. 2, pp. 99–127, 2002.
- [39] E. Real *et al.*, "Large-scale evolution of image classifiers," in *Proc. ICML*, 2017, pp. 2902–2911.
- [40] S. L. Lauritzen and D. J. Spiegelhalter, "Local computations with probabilities on graphical structures and their application to expert systems," *J. Roy. Stat. Soc. B*, vol. 50, no. 2, pp. 157–224, 1988.
- [41] I. A. Beinlich, H. J. Suermondt, R. M. Chavez, and G. F. Cooper, "The ALARM monitoring system: A case study with two probabilistic inference techniques," in *Proc. AIME*, 1989, pp. 247–256.
- [42] Z. Liu, P. Luo, X. Wang, and X. Tang, "Deep learning face attributes in the wild," in *Proc. ICCV*, 2015, pp. 3730–3738.
- [43] J. du Terrail *et al.*, "FLamby: Datasets and benchmarks for cross-silo federated learning in realistic healthcare settings," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2022.

Project Guide: _____

Students: _____