

Causal Proof of Reasoning: Defending Federated Learning Against Explanation Poisoning via Structural Auditing

Ayush Chadha, Ragav Palaniswamy, Sanjana Gummuluru, and Sridevi M

Abstract—Federated Learning (FL) enables collaborative model training across distributed, privacy-sensitive participants without centralizing raw data. However, its decentralized trust model creates two severe vulnerability classes: *backdoor poisoning*, where malicious clients embed hidden triggers into the global model, and *Explanation Poisoning*, a sophisticated dual-objective attack in which an adversary simultaneously optimizes for malicious inference and benign-looking post-hoc explanations, rendering LIME, SHAP, and gradient-based auditing tools dangerously misleading. Existing Byzantine-defense aggregators analyze weight-vector divergence and therefore fail against both attack classes: modern attacks exploit the overparameterization of neural networks to preserve statistical similarity while corrupting latent dependency structure.

We propose Causal Proof of Reasoning (PoR), a structural auditing framework that shifts the defense layer from numerical weight analysis to topological inspection of latent feature dependencies. Each client applies the NOTEARS algorithm to recover a latent dependency structure under acyclicity constraints, producing a Directed Acyclic Graph (DAG) as a structural certificate. A server-side SimGNN Logic Validator compares each DAG against a global consensus, approximating Graph Edit Distance (GED) in $\mathcal{O}(1)$ time. Poisoned clients exhibit detectable structural deformation that cannot be concealed without simultaneously destroying the backdoor, a topological irreducibility that forms the core defense principle.

PoR is evaluated within a heterogeneous neuroevolutionary federated framework (FedNEAT) and validated on two Bayesian network benchmarks and an Algorithmic Finance environment. Evaluation across 5 independent runs shows that PoR reduces Attack Success Rate from 60.30% to 0.40% on ALARM (at 38.75% MTA; see Section VII for a discussion of this utility gap) and from 5.62% to 0.52% on ASIA (at 93.95% MTA), with zero observed false positives on the Bayesian network benchmarks. Unlike parameter-space Byzantine defenses, PoR audits latent dependency structures, shifting federated security from weight divergence analysis to representation-structure verification.

Index Terms—Federated learning, explanation poisoning, causal discovery, graph neural networks, Byzantine robustness, structural auditing.

I. INTRODUCTION

The authors are with the Department of Computer Science and Engineering, National Institute of Technology, Tiruchirappalli, Tamil Nadu 620015, India (e-mail: {106122027, 106122095, 106122106, msridevi}@nitt.edu).

FEDERATED Learning (FL) [1] has emerged as a foundational paradigm for training machine learning models across distributed, data-heterogeneous participants without centralizing raw data. By aggregating only model updates rather than datasets, FL offers strong privacy guarantees and enables collaboration across health-care, finance, and autonomous systems, domains where data governance prohibits centralized collection [2], [3]. Recent advances in federated system design increasingly require participants to operate on heterogeneous data distributions and evolving local objectives, introducing structural opacity: a participant’s learned internal representations may diverge measurably from the honest population without triggering parameter-level statistical alarms.

Despite these advantages, FL’s distributed trust model exposes it to a fundamental class of threats: *model poisoning and backdoor attacks*. A malicious participant can submit crafted model updates that embed a hidden trigger into the global model, causing it to misclassify any input containing that trigger while maintaining normal accuracy on clean inputs [4], [5]. Unlike data poisoning, model poisoning attacks operate entirely within the update protocol—the adversary’s submission is indistinguishable from legitimate updates at the parameter level. More sophisticated Distributed Backdoor Attacks (DBA) [6] decompose a global trigger across multiple malicious agents, each submitting a seemingly benign partial pattern that assembles into a potent trigger only at the global model, yielding Attack Success Rates (ASR) exceeding 85% even against state-of-the-art defense aggregators.

A. The Insufficiency of Weight-Space Defenses

Existing Byzantine-resilient defenses assume that adversarial updates measurably diverge in parameter space. However, modern backdoor and explanation-poisoning attacks exploit the overparameterization of neural networks to preserve statistical similarity while corrupting latent reasoning structure. This creates a fundamental evasion surface that weight-space auditing cannot close.

The dominant family of such aggregators, including Krum [7], Trimmed Mean, and Median [8]—operates by computing pairwise distances or statistics over the high-dimensional weight vectors submitted by clients, discarding statistical outliers. These methods embed a critical assumption: *a malicious update must be statistically*

anomalous in weight space. This assumption is routinely violated.

An adversary who scales their poisoned update to match the ℓ_2 -norm of honest updates, or who mounts a DBA with sub-triggers that individually appear within the honest distribution, will evade all weight-divergence detectors. Even FLTrust [9], which bootstraps trust from a server-held root dataset, is susceptible when the trigger is semantically aligned with legitimate feature correlations. A parallel vulnerability is exploited at the explainability layer, a class of attacks we term *Explanation Poisoning* [10], [11]: the adversary solves a dual-objective optimization problem

$$\theta^* = \arg \min_{\theta} \underbrace{\left(\mathcal{L}_{\text{CE}}(f_{\theta}(x_{\text{trig}}), y_{\text{tgt}}) \right)}_{\text{Backdoor success}} + \beta \cdot \underbrace{\left\| E_{\theta}(x_{\text{trig}}) - E_{\theta_{\text{clean}}}(x_{\text{clean}}) \right\|_2^2}_{\text{Explanation camouflage}} \quad (1)$$

where $E_{\theta}(x)$ is a post-hoc explanation function (e.g. SHAP, GradCAM). The model effectively *lies* to the human auditor: it makes decisions based on a hidden trigger while generating explanations that highlight legitimate features, satisfying post-hoc auditing tools such as LIME and SHAP while remaining malicious in deployment [12].

The core failure mode is conceptual: *these defenses audit what a model predicts under specific inputs, not how it reasons internally*. An adversary who preserves the statistical signature of their weight updates while corrupting the underlying reasoning structure will evade detection indefinitely.

B. Causal Proof of Reasoning

We argue that defending federated systems requires auditing not merely what models predict, but how internal latent dependencies are structurally organized. Rather than inspecting weight magnitudes or output distributions, we inspect the *latent structural dependency* of a client’s learned representations, the topology of relationships between latent concepts that organizes internal feature processing. A model that has been poisoned must corrupt its internal feature associations to link the trigger with the target class; this corruption is reflected as a topological deformation of the latent Directed Acyclic Graph (DAG) relative to the honest global consensus. Crucially, Equation (1)’s explanation-camouflage term can falsify output-space attributions but cannot hide the structural deformation of the latent dependency graph without simultaneously destroying the backdoor itself, a topological irreducibility that forms the core defense principle.

Our framework, **Causal Proof of Reasoning (PoR)**, operationalizes this intuition as follows. Each client applies the NOTEARS algorithm—a continuous optimization procedure that enforces strict DAG constraints [13]—to its latent representation layer to recover a latent dependency structure under acyclicity constraints, producing a compact DAG as a structural certificate. The server maintains a *Global Consensus Graph*—a momentum-updated

aggregate of past honest client structures, and employs a pre-trained SIMGNN [14] to predict the Graph Edit Distance (GED) between each submitted DAG and the consensus in $\mathcal{O}(1)$ inference time, regardless of graph size. A complementary Causal Effect Distance (CED) gate audits the NOTEARS coefficient matrix to catch weight-only backdoors that preserve graph topology. Clients whose structural certificates deviate measurably from the consensus are identified as adversarial and excluded from aggregation.

PoR is evaluated within FEDNEAT—a Federated NeuroEvolution of Augmenting Topologies framework where clients evolve heterogeneous neural architectures via genetic mutation, eliminating the homogeneous gradient signature that statistical defenses rely on and making it the most adversarially challenging substrate for structural auditing. We further validate PoR in an Algorithmic Finance environment as validation beyond static Bayesian networks, demonstrating generalizability to sequential decision processes. The system is implemented using the Flower federated learning framework [15] with Ray-based parallelism.

C. Contributions

The primary contributions of this work are:

- 1) **Causal Proof of Reasoning (PoR)**: A structural latent dependency auditing defense for FL that recovers DAGs from client latent spaces via NOTEARS and validates them via SIMGNN, operating orthogonally to weight-space defenses and defeating Explanation Poisoning by exploiting the topological irreducibility of backdoor execution.
- 2) **Latent Deformation Evidence**: Empirical evidence via Wasserstein distance ($\bar{W}_1 = 5.25$) that backdoor poisoning induces measurable latent dependency deformation, providing the theoretical justification for structural auditing as a security primitive in federated systems.
- 3) **Dynamic Percentile Thresholding**: An adaptive rejection boundary derived from the expected Byzantine ratio, preventing honest client rejection under non-IID data drift across diverse graph complexity regimes.
- 4) **Empirical Validation**: Comprehensive evaluation on ASIA and ALARM Bayesian networks (ASR suppression from 60.30% to 0.40% on ALARM and from 5.62% to 0.52% on ASIA), extended with a Finance agentic environment and dual-gate logic validation (GED + CED), demonstrating generalizability across graph complexity regimes and sequential decision domains.

D. Agentic Federated Learning and the Trust Gap

The emergence of agentic federations—where clients curate their own training data, optimize local reward

functions, and operate in dynamic environments such as healthcare cross-silo networks [16] and biometric analysis systems [17], introduces a deeper vulnerability. An agentic node can silently alter its learning process without triggering per-round statistical alarms, a threat we term *Memory Poisoning*. Regulatory frameworks such as the EU AI Act mandate explainability for high-risk AI deployments, but they do not require explanation integrity. An adversary who falsifies saliency maps while maintaining accuracy complies with the letter of the regulation while violating its spirit (transforming regulatory compliance itself into an attack surface). PoR closes this trust gap by enforcing structural alignment of causal reasoning rather than auditing output-space explanations.

The remainder of this paper is organized as follows. Section II surveys related work. Section III formalizes the federated system and threat model. Section IV presents the PoR methodology. Section V describes the experimental setup. Section VI presents results and analysis. Section VII discusses implications and limitations. Section VIII concludes.

II. RELATED WORK

The literature on federated learning security spans three overlapping threads: statistical anomaly detection in parameter space, causal structure learning for feature engineering, and neural graph distance metrics for graph matching. We survey each thread, identify where existing approaches fall short against explanation-aware adversaries, and contextualize the precise gap that PoR addresses.

A. Byzantine-Robust Federated Aggregation

The foundational challenge of FL under Byzantine faults was formalized by Blanchard et al. [7], who proposed *Krum*—a server-side aggregator that selects the client update with the minimum sum of distances to its nearest neighbours, provably tolerating up to $f < n/2$ Byzantine clients. Yin et al. [8] established statistical optimality bounds for coordinate-wise Trimmed Mean and Median aggregators. Despite theoretical guarantees, these methods share the assumption that adversarial updates are statistical outliers in weight space, an assumption systematically violated by norm-constrained and DBA-style attacks.

FLTrust [9] addresses this by maintaining a server-held root dataset to compute trust scores for each client’s update direction, gating aggregation by cosine similarity. Our baseline in this work—*Byzantine-Robust FedAvg*—implements a similar cosine weight-divergence filter, representing the current practical state of the art for deployed systems. FLAME [18] combines clustering with adaptive noise injection but introduces significant utility degradation. Recent reconstruction-error-based approaches such as FedREDefense [19] (ICML 2024) detect anomalous updates via autoencoder reconstruction error, but fail against Explanation Poisoning attacks, which are deliberately

clean and structured. ESFL [20] (IEEE TDC 2025) employs adaptive local differential privacy and compressive sensing for security, but the injected noise paradoxically aids attackers by obscuring malicious triggers. All of these approaches operate entirely in parameter space and cannot inspect the *internal latent dependency structure* of client models.

B. Backdoor Attacks and Explanation Poisoning

Backdoor attacks in FL were systematically studied by Bagdasaryan et al. [4], who showed that a single malicious client can dominate the global model’s behavior through *model replacement*—scaling their poisoned update by $1/q$ where q is the learning rate. Chen et al. [5] demonstrated trigger-based targeted misclassification. Xie et al. [6] introduced Distributed Backdoor Attacks (DBA), decomposing a global trigger across multiple participants to evade per-client anomaly detection. Sun et al. [21] showed that triggers persist across aggregation rounds; Wang et al. [22] demonstrated tail-distribution backdoors that survive standard aggregation. A critical property shared across these attacks is that they exploit overparameterization to preserve statistical camouflage in weight space while corrupting internal representations.

A parallel attack surface exploits post-hoc explainability tools. Liyanage et al. [10] demonstrated Scaffolding Attacks in federated intrusion detection, where an adversary embeds a persistent bias while concealing it from SHAP. Slack et al. [11] showed that LIME and SHAP explanations can be systematically falsified via an Out-of-Distribution (OOD) switching mechanism. Heo et al. [12] demonstrated the same for gradient-based saliency maps. These attacks formalize as the dual-objective of Equation (1): the adversary preserves statistical camouflage in both weight space and output-space explanations simultaneously. PoR defends against this by auditing latent structural dependencies rather than output explanations, rendering the statistical camouflage mechanism irrelevant at the representation layer.

C. Causal Discovery and Graph Auditing

Pearl’s framework [23] established causal DAGs as the canonical language for encoding conditional independence relationships. Peters et al. [24] provided a rigorous treatment of identifiability conditions for causal structures from observational data. The NOTEARS algorithm [13] revolutionized computationally tractable causal structure learning by reformulating the combinatorial DAG constraint as a differentiable equality constraint. We apply NOTEARS to latent representations rather than raw features, treating the learned representation space as the observational domain for latent dependency structure recovery.

Exact GED computation is NP-hard [25]. SIMGNN [14] introduced a Siamese GNN architecture that learns to predict normalized GED directly, achieving $\approx 1720\times$ speedup over exact A* search. GRAPHEDX [26] (NeurIPS 2024)

confirmed the feasibility of real-time graph-based governance at federation scale. Unlike static graph distances such as Jaccard similarity or exact GED, SIMGNN learns adaptive semantic similarity under evolving consensus graphs—a property essential in federated settings where honest graph topologies naturally shift across rounds.

Federated Causal Discovery (FCD) has emerged as a parallel direction: FedECD [27] (NeurIPS 2024) and Fed-Causal [28] (AAAI 2025) both use causal graphs to improve accuracy under non-IID data. Importantly, none of these approaches consider adversarial clients intentionally submitting deformed latent dependency graphs, nor use latent topology as a federated security gate. This is the precise gap PoR addresses.

D. *NeuroEvolution in Federated Settings*

Stanley and Miikkulainen’s NEAT algorithm [29] introduced topology-evolving neural networks where both weights and connections evolve via genetic operators, producing heterogeneous architectures from a minimal initial structure. FEDNEAT extends this to the federated setting, creating heterogeneous client architectures without homogeneous weight vectors, rendering standard cosine-distance defenses inapplicable and making FEDNEAT the most stringent testbed for PoR’s topology-based auditing.

E. *Known Limitations of Causal Discovery and Interpretability*

The NOTEARS algorithm [13] is computationally tractable but carries strict assumptions: it optimally recovers linear-Gaussian DAGs, and in non-linear or high-dimensional regimes, recovered graphs may contain spurious edges from algorithmic artifacts rather than adversarial intent. Furthermore, NOTEARS lacks scale-invariance; a simple linear rescaling of a feature column can alter the inferred DAG even from structurally identical data. In non-IID federated settings, honest clients on different data pipelines may submit divergent graphs due to normalization differences, not malicious behavior. PoR addresses this by calibrating the rejection threshold τ using the empirical distribution of GED scores from observed honest clients rather than a fixed domain-agnostic value. Future iterations may integrate interpretable graph matchers such as GEDGNN [30] to provide auditable rejection justifications.

F. *Generational Analysis of Federated Defenses*

Federated defenses have evolved through three generations, each progressively closer to but still fundamentally incapable of detecting Explanation Poisoning. First-generation statistical outlier detectors (Krum [7], Trimmed Mean [8]) assume adversarial updates are Euclidean outliers, an assumption directly violated by norm-constrained and DBA-style attacks. Second-generation similarity and trust approaches (Fools-Gold [31], FLTrust [9]) fail against Explanation Poisoning because honest agents in heterogeneous domains

TABLE I
RESEARCH GAP ADDRESSED BY PoR.

Domain	State-of-the-art technique or Existing approach (2024–2026)	PoR Contribution
Attack	Explanation Poisoning undetectable by SHAP [10], [11]	Validates latent structure, not output explanations
Defense	Weight statistics or DP [19], [20]	Structural logic-based auditing gate
Causality	Used for accuracy [27], [28]	Latent topology as a security primitive
GED	Neural GED solved [26] but unused in FL	Neural GED as real-time security gatekeeper

produce similar updates, triggering false Sybil flags, and Root Datasets may not cover rare adversarial conditions. Third-generation reconstruction-error methods (FedREDefense [19], FLAME [18]) detect anomalous weights but pass Explanation Poisoning attacks, which are deliberately clean and structured. PoR operates orthogonally to all three generations by auditing causal logic graphs rather than weight vectors, rendering the attack’s statistical camouflage mechanism irrelevant.

G. *Summary of the Research Gap*

Table I contextualizes PoR against the current state of the art across four critical dimensions.

III. PROBLEM FORMULATION AND THREAT MODEL

A. *Federated Learning System Model*

Consider a federation of N clients $\mathcal{C} = \{c_1, c_2, \dots, c_N\}$ coordinated by a central server. Each client c_i holds a private local dataset $\mathcal{D}_i$ drawn from a distribution P_i over $\mathcal{X} \times \mathcal{Y}$, where $\mathcal{X} \subseteq \mathbb{R}^d$ is the feature space and $\mathcal{Y} = \{1, \dots, K\}$ is the label space. Datasets are non-IID: $P_i \neq P_j$ in general. The federation proceeds over T global rounds; in each round t , each client trains a local model on $\mathcal{D}_i$ and submits a structured update $U_i^{(t)}$ to the server, which aggregates across clients and broadcasts an updated global model $\theta^{(t)}$.

In FEDNEAT, updates $U_i^{(t)}$ are evolved genome structures rather than gradient vectors. Under PoR, each update additionally includes a structural certificate $\mathcal{G}_i^{(t)}$ —a DAG over the latent feature space.

B. *Threat Model*

We model a *Byzantine adversary* controlling a subset $\mathcal{B} \subset \mathcal{C}$ of $|\mathcal{B}| = m$ malicious clients, where $m < N/2$. The adversary has:

- **Local omniscience:** Full control over $\mathcal{D}_i$ for $c_i \in \mathcal{B}$, including data poisoning and model manipulation.
- **Protocol compliance:** The adversary submits updates that conform to the FL protocol. They do *not* compromise the server or honest clients’ local data.

- **Trigger knowledge:** The adversary coordinates a shared backdoor trigger $\delta \in \mathcal{X}$ (or distributes sub-triggers across $\mathcal{B}$ as in DBA [6]).

Definition 1 (Backdoor Attack). *A backdoor attack succeeds if the global model $\theta^{(T)}$ satisfies:*

$$f_{\theta^{(T)}}(x) = y \quad \forall (x, y) \in \mathcal{D}_{clean} \quad (2)$$

$$f_{\theta^{(T)}}(x \oplus \delta) = y_t \quad \forall x \in \mathcal{X} \quad (3)$$

where $y_t \neq y$ is the adversary's target label and $\oplus$ denotes trigger injection into feature δ .

C. Poisoning Strategy

Each malicious client $c_i \in \mathcal{B}$ poisons a fraction ρ of its local training data by (i) setting a designated trigger feature $x_{[\delta]} \leftarrow 0$ and (ii) flipping the corresponding label to the target class y_t . This strategy is designed to maximize the probability that the model's latent representations associate the trigger feature with y_t while preserving accuracy on clean samples.

D. Defense Objective

A defense mechanism Φ is *effective* if, given a federation with m Byzantine clients, the protected global model $\theta_{\Phi}^{(T)}$ satisfies:

$$\text{ASR}(\theta_{\Phi}^{(T)}) \ll \text{ASR}(\theta_{\text{undefended}}^{(T)}) \quad (4)$$

$$\text{MTA}(\theta_{\Phi}^{(T)}) \approx \text{MTA}(\theta_{\text{honest}}^{(T)}) \quad (5)$$

PoR aims to achieve this by replacing weight-space filtering with structural certificate validation.

IV. PROPOSED METHODOLOGY

Fig. 1 illustrates the end-to-end PoR pipeline. Each federated round proceeds in four stages: (i) local genome evolution and feature extraction by clients, (ii) latent causal graph extraction via NOTEARS, (iii) server-side structural auditing via SIMGNN, and (iv) conditional aggregation of accepted updates only.

A. FedNEAT: Heterogeneous Client Architecture

Unlike conventional FL, where all clients train identical architectures via gradient descent, FEDNEAT assigns each client a private *genome* Γ_i —a graph-encoded neural network topology that evolves through mutation across rounds. Following the NEAT formalism [29], each genome is a directed graph $\Gamma_i = (V_i, E_i)$ where nodes V_i represent neurons and edges E_i represent weighted synaptic connections tagged with globally unique *innovation numbers*.

In each round, each client c_i performs L local epochs of forward propagation on $\mathcal{D}_i$ using its current genome, computes a fitness score based on classification loss, and applies NEAT mutation operators: weight perturbation (probability $p_w = 0.8$), connection addition ($p_c = 0.1$), and node insertion ($p_n = 0.05$). The evolved genome is serialized and submitted to the server as the structural update $U_i^{(t)}$.




Fig. 1. The PoR federated round. Clients extract latent DAGs via NOTEARS and submit them alongside genome updates. The server Logic Validator (SIMGNN) computes structural deviation from the Global Consensus Graph and rejects adversarial submissions before aggregation.

Importantly, because genomes are heterogeneous—different clients may have different numbers of nodes and connections—there is no canonical weight-vector representation. This renders standard defenses based on cosine distance and Euclidean norms completely inapplicable, establishing FEDNEAT as the most stringent possible testbed for PoR.

B. Client-Side Architecture Overview

Each client in the PoR framework operates as a deliberative agent with three functional modules. The *Perception Module* implements causal disentanglement, separating invariant causal factors from spurious correlations in the input representation. The *Cognitive Module* extracts the causal logic DAG via NOTEARS applied to the latent representation layer, converting the learned dependency structure into a structural certificate $\mathcal{G}_i$. The *Action Module* assembles the structured update payload $\{\Delta\theta_i, \mathcal{G}_i, W_i\}$ transmitted to the server, where $\Delta\theta_i$ is the weight update, $\mathcal{G}_i$ is the causal graph, and W_i is the NOTEARS coefficient matrix. This tri-module design ensures that each client transmits not only its performance signal but also a verifiable reasoning certificate, enabling server-side structural auditing independent of model architecture.

C. Latent Dependency Structure Recovery via NOTEARS

Backdoor optimization deforms latent manifolds by introducing unnatural feature-target couplings. These deformations alter conditional dependency structure among latent features, forcing NOTEARS to recover structurally divergent graphs. The structural certificate is therefore sensitive to poisoning in proportion to the severity of the backdoor's latent corruption.

After local training, each client recovers a latent dependency structure under acyclicity constraints from its learned latent representations. Let $Z_i \in \mathbb{R}^{n_i \times p}$ denote the matrix of p -dimensional latent activations collected over a local mini-batch of size n_i . We treat each column of

Z_i as a random variable and seek a weighted adjacency matrix $W_i \in \mathbb{R}^{p \times p}$ encoding the conditional independence structure.

1) *DAG Constraint*: The NOTEARS algorithm [13] casts DAG structure learning as a continuous optimization problem. The key insight is the algebraic characterization of acyclicity:

Theorem 1 (Acyclicity Constraint [13]). *A matrix $W \in \mathbb{R}^{p \times p}$ is a DAG (i.e., has no directed cycles) if and only if*

$$h(W) = \text{tr}(e^{W \circ W}) - p = 0 \quad (6)$$

where $\circ$ denotes the Hadamard (element-wise) product.

Each client solves the constrained optimization:

$$\min_{W_i} \frac{1}{2n_i} \|Z_i - Z_i W_i\|_F^2 + \lambda \|W_i\|_1 \quad \text{s.t.} \quad h(W_i) = 0 \quad (7)$$

where the ℓ_1 penalty with coefficient λ enforces sparsity, promoting interpretable DAGs with few edges. Problem (7) is solved via augmented Lagrangian optimization [13] for a maximum of $I_{\max}$ iterations.

2) *Graph Construction*: The resulting adjacency matrix W_i is thresholded at ϵ : edge (u, v) is retained if $|W_i[u, v]| > \epsilon$. The threshold ϵ is a global hyperparameter controlling DAG sparsity (default $\epsilon = 0.02$, tuned per dataset). The extracted DAG $\mathcal{G}_i = (V, E_i)$ shares a fixed node set V corresponding to the p latent dimensions and has a dataset-specific edge set E_i .

D. Latent Distribution Shift Analysis

To provide theoretical grounding for the structural auditing approach, we analyze how backdoor poisoning affects the latent feature distributions.

Proposition 1 (Latent Deformation Under Poisoning). *Let $Z^{(h)}$ and $Z^{(\text{adv})}$ denote the latent representations of an honest and an adversarial client respectively. Under a feature-zeroing backdoor attack with poison fraction ρ , the adversarial latent distribution is deformed with expected Wasserstein-1 distance:*

$$\mathbb{E}[W_1(Z^{(h)}, Z^{(\text{adv})})] > 0 \quad (8)$$

as the trigger feature induces a systematic shift in the joint feature distribution that propagates through the latent layer.

We empirically measure this deformation by running matched honest and adversarial clients on the ASIA dataset and computing the per-dimension Wasserstein-1 distance between their latent activation distributions. The mean Wasserstein distance across all p latent dimensions is $\bar{W}_1 = 5.2542$, confirming that poisoned clients produce structurally distinct latent manifolds. This result validates the core premise of PoR: structural deformation at the latent level is both *necessary* and *detectable* as a consequence of the poisoning strategy.

E. SimGNN Logic Validator

The Logic Validator $\mathcal{V}$ computes a structural divergence score for each submitted DAG $\mathcal{G}_i$ relative to the server-maintained Global Consensus Graph $\mathcal{G}^*$.

1) *Architecture*: We adopt the SIMGNN architecture [14], a Siamese GNN that encodes two graphs independently and predicts their normalized GED from the resulting embeddings. Each branch applies L_{conv} layers of graph convolution:

$$h_v^{(l+1)} = \sigma\left(W^{(l)} \cdot \text{MEAN}\left(\{h_u^{(l)} : u \in \mathcal{N}(v) \cup \{v\}\}\right)\right) \quad (9)$$

where $\mathcal{N}(v)$ denotes the neighbourhood of node v , $W^{(l)} \in \mathbb{R}^{d_l \times d_{l-1}}$ are learnable weights, and σ is the ReLU activation. After L_{conv} layers, a global graph embedding is obtained via dual pooling:

$$\mathbf{g} = [\text{MEAN}(\{h_v^{(L)}\}) \parallel \text{MAX}(\{h_v^{(L)}\})] \in \mathbb{R}^{2d_L} \quad (10)$$

The concatenated embeddings $[\mathbf{g}_i \parallel \mathbf{g}^*]$ from both branches are passed through a two-layer MLP regression head to predict the normalized GED score $\hat{s}_i \in [0, 1]$.

2) *Node Feature Encoding*: Since Bayesian network nodes are categorical variables with dataset-specific semantics, we use padded one-hot encodings of length $p_{\max}$ (set to the maximum node count across all supported datasets). For a graph with p nodes, each node v_j has feature $\mathbf{x}_j = \mathbf{e}_j \in \{0, 1\}^{p_{\max}}$, where $\mathbf{e}_j$ is the j -th standard basis vector. Zero-padding ensures a fixed input dimensionality across different graph sizes.

3) *Pre-training via Curriculum*: The Logic Validator is pre-trained on the server before the federation begins, using the Global Consensus Graph as the anchor. Training pairs $(\mathcal{G}^*, \mathcal{G}^{\text{mut}})$ are generated by randomly mutating $\mathcal{G}^*$ (adding, removing, or reversing edges with probability $p_{\text{div}} = 0.1$ per edge) and computing the exact GED of small graphs (≤ 10 nodes) or an upper-bound approximation for larger graphs. The model is trained to minimize the mean squared error between the predicted normalized GED $\hat{s}_b$ and the ground-truth (or approximated) normalized GED $\tilde{s}_b$:

$$\mathcal{L}_{\text{SIMGNN}} = \frac{1}{B} \sum_{b=1}^B (\hat{s}_b - \tilde{s}_b)^2 \quad (11)$$

Pre-training converges in $E = 500$ epochs with Adam optimizer ($\eta = 0.001$, batch size $B = 32$).

Unlike static graph distances such as Jaccard similarity or exact GED, SIMGNN learns adaptive semantic similarity under evolving consensus graphs—making it resilient to naturally shifting honest-client topology distributions across federation rounds.

F. Causal Effect Distance (CED) Gate

(Implementation complete; standalone ablation benchmarking is deferred to future work. CED can be isolated by running PoR with GED disabled and measuring ASR against weight-only backdoors that preserve topology while scaling edge coefficients.)

While SIMGNN detects topological deformation, an adversary may attempt a *weight-only backdoor* by maintaining the exact consensus topology but measurably altering

NOTEARS coefficient magnitudes. The CED Gate counters this by computing the maximum absolute divergence over consensus edges:

$$\text{CED}_i = \max_{(u,v) \in E^*} |W_i[u,v] - W^*[u,v]| \quad (12)$$

Clients are rejected via MAD filtering if their CED score is a statistical outlier ($\gamma = 3.0$). Together, the GED and CED gates form a dual-barrier logic validator catching both structural and magnitude-based backdoor variants.

G. Dynamic Percentile Thresholding

A critical challenge for any anomaly detector in FL is distinguishing true adversaries from honest clients who exhibit atypical structural updates due to non-IID data heterogeneity. A static threshold τ is brittle: too low and honest stragglers are rejected; too high and adversaries pass.

We address this with a *dynamic rank-preserving percentile threshold*. In each round t , the Logic Validator computes predicted GED scores $\{\hat{s}_i\}_{i=1}^N$ for all N submitted DAGs. Given an expected Byzantine fraction $\beta = m/N$, the honest ratio is $\alpha = 1 - \beta = (N - m)/N$. The dynamic rejection boundary is:

$$\tau^{(t)} = \text{Percentile}(\{\hat{s}_i\}, \max(50, \alpha \cdot 100 \cdot 0.95)) \quad (13)$$

Any client c_i with $\hat{s}_i > \tau^{(t)}$ is flagged as adversarial and excluded from aggregation. The factor 0.95 provides a small safety margin below the honest boundary, ensuring near-zero false positive rates even when honest clients experience transient structural drift.

An absolute floor $\tau_{\min}$ (dataset-dependent; 0.35 for ALARM, 0.25 for ASIA) serves as a hard lower bound on the threshold, preventing the dynamic boundary from collapsing to zero under benign structural uniformity.

H. PoR-Augmented Aggregation Protocol

Algorithm 1 formalizes the complete server-side aggregation procedure. The accepted clients' genome updates are merged using innovation-number-aligned weighted averaging. To prevent collusion attacks where multiple adversaries slowly drift the consensus graph over time, the Global Consensus Graph is updated using a Bayesian Dirichlet posterior update rather than simple Exponential Moving Average.

Each potential edge $e = (u, v)$ in the consensus maintains Beta parameters (α_e, β_e) tracking its credibility:

$$\alpha_e^{(t+1)} = \alpha_e^{(t)} + \sum_{i \in \mathcal{A}^{(t)}} \mathbb{I}[e \in \mathcal{G}_i^{(t)}] \quad (14)$$

$$\beta_e^{(t+1)} = \beta_e^{(t)} + \sum_{i \in \mathcal{A}^{(t)}} \mathbb{I}[e \notin \mathcal{G}_i^{(t)}] \quad (15)$$

An edge is retained in the global consensus topology only if its Bayesian credibility exceeds the consensus momentum threshold $\mu = 0.70$:

$$\frac{\alpha_e^{(t+1)}}{\alpha_e^{(t+1)} + \beta_e^{(t+1)}} > \mu \quad (16)$$

Algorithm 1: PoR-Augmented Federated Aggregation

Input: Global consensus $\mathcal{G}^*$, weights W^* , pre-trained $\mathcal{V}$, client submissions $\{(U_i, \mathcal{G}_i, W_i)\}_{i=1}^N$, Byzantine fraction m/N , momentum μ
Output: Updated $\theta^{(t+1)}$, updated $\mathcal{G}^{*(t+1)}$

// Stage 1: Score submitted graphs (GED gate)

1 **for** $i \leftarrow 1$ **to** N **do**
2 | $\hat{s}_i \leftarrow \mathcal{V}(\mathcal{G}_i, \mathcal{G}^*)$ // GED score
3 **end**

// Stage 2: Dynamic GED threshold

4 $\alpha \leftarrow (1 - m/N) \cdot 0.95$
5 $\tau \leftarrow \text{Percentile}(\{\hat{s}_i\}, \max(50, \alpha \cdot 100))$

// Stage 3: Accept clients below threshold

6 $\mathcal{A} \leftarrow \{i \mid \hat{s}_i \leq \tau\}$

// Stage 4: Aggregate accepted updates

7 $\theta^{(t+1)} \leftarrow \text{FedAgg}(\{U_i\}_{i \in \mathcal{A}})$

// Stage 5: Bayesian Consensus Update

8 Update posteriors (α_e, β_e) via accepted edges

9 $\mathcal{G}^{*(t+1)} \leftarrow \{e \mid \alpha_e / (\alpha_e + \beta_e) > \mu\}$

// Stage 6: Fine-tune Logic Validator

10 Fine-tune $\mathcal{V}$ on $\{(\mathcal{G}^*, \mathcal{G}_i)\}_{i \in \mathcal{A}}$

11 **return** $\theta^{(t+1)}$, $\mathcal{G}^{*(t+1)}$

This mechanism ensures edges seen strongly in past rounds survive single-round fluctuations (e.g., due to curriculum learning transitions), while requiring sustained, overwhelming evidence to introduce new adversarial structures.

I. Differential Privacy via DP-SGD

(Implementation complete in the Finance environment; formal ablation benchmarking—measuring MTA/ASR vs. privacy budget ϵ —is deferred to future work.)

PoR implements Differential Privacy via DP-SGD [32] at the client level prior to DAG extraction, clipping per-sample gradients to ℓ_2 -norm bound C and injecting Gaussian noise $\mathcal{N}(0, \sigma^2 C^2)$. Privacy cost across T rounds is tracked via Rényi Differential Privacy (RDP) composition [33] (see Algorithm 1, Stage 4). PoR's adversary filtering improves the privacy-utility trade-off: by rejecting adversarial gradients prior to aggregation, the effective noise impact on honest consensus topology is reduced.

J. Market Simulator for the Finance Domain

To validate PoR under dynamic agentic conditions, we construct a correlated market simulator across $M = 11$ GICS sectors. Three mathematical components govern the price dynamics.

1) *Ornstein-Uhlenbeck Mean Reversion*: Sector prices exhibit mean reversion modeled as a discretized Ornstein-Uhlenbeck process [34]:

$$\Delta p_{i,t} = \mu_R + \kappa \left(\frac{b_i - p_{i,t}}{b_i} \right) + \sigma_{i,t} Z_{i,t} \quad (17)$$

where μ_R is the regime-specific drift under a Markov chain (Bear/Sideways/Bull), $\kappa = 0.001$ is the mean-reversion speed, and $\sigma_{i,t}$ is the GARCH-conditional volatility.

2) *GARCH(1,1) Volatility Clustering*: Conditional variance follows GARCH(1,1) [35]:

$$\sigma_{i,t}^2 = \omega + \alpha r_{i,t-1}^2 + \beta \sigma_{i,t-1}^2 \quad (18)$$

where $r_{i,t-1}$ is the previous log return, introducing non-stationarity that continuously stress-tests the federated defense.

3) *Cross-Sector Correlation*: Correlated shocks are generated via Cholesky decomposition $\Sigma = LL^\top$ of an empirical inter-sector correlation matrix: $Z_t = L\epsilon_t$, $\epsilon_t \sim \mathcal{N}(0, I)$. Adversarial agents attempt to exploit a VIX macro-trigger injected on top of these shocks, and are detected via the structural GED gate as their causal query graphs diverge from the honest consensus topology.

V. EXPERIMENTAL SETUP

A. Datasets

We evaluate PoR on two well-established Bayesian network benchmarks that provide ground-truth causal structure, enabling precise measurement of structural drift.

ASIA [36]: An 8-node, 8-arc Bayesian network modelling lung disease diagnosis (Asia travel, tuberculosis, lung cancer, bronchitis, dyspnea, X-ray). Each node is a binary variable. We generate 10,000 samples using ancestral sampling via `pgmpy` [37]. The attack targets the *lung* variable (binary, 2 classes) with trigger feature index 0 and target label 0, using poison fraction $\rho = 0.20$.

ALARM [38]: A 37-node, 46-arc Bayesian network modelling anaesthesia monitoring. Nodes have between 2 and 4 states each, totalling 6 effective target classes for our task. We generate 10,000 samples and attack the *bp* (blood pressure) variable with trigger feature index 0 and target label 2, using a stronger poison fraction $\rho = 0.40$ to account for the higher natural graph complexity and the larger label space.

ALARM is the primary static security benchmark: its 37-node topology introduces substantially higher structural variance among honest clients and requires a potent attack to overcome the natural causal complexity. ASIA serves as a dense-graph control benchmark to validate generalization across graph density regimes.

Algorithmic Finance (Hedge Fund): To validate PoR in an agentic environment with dynamic, sequential decision making, we implement a market trading simulation. The environment state is a 70-dimensional continuous vector (33 market API visibility masks, 36 revealed API values, 1 VIX macro trigger). Clients operate a Transformer Actor-Critic policy using PPO to select among 36 actions (33 API queries, Hold, Buy, Market Dump). The topology here is not a static BN, but a *Cognitive Execution Graph* extracted dynamically from the agent’s action sequences via NOTEARS. This domain enforces a Curriculum Learning schedule, scaling from 5 to 11 active sectors, creating immense variance in honest topologies across early rounds.

B. Federated Configuration

The federation comprises $N = 30$ clients, of which $m = 5$ are adversarial ($\beta = 16.7\%$ Byzantine fraction). Each round, all 30 clients participate. Local training runs for $L = 5$ epochs per round, and the federation runs for $T = 15$ global rounds—sufficient for NEAT genomes to converge on the ASIA search space and sufficient to show sustained attack suppression on ALARM.

Client parallelism is managed by Ray [39] with 4 CPU cores per actor. Experiments were conducted on a cloud compute instance with multi-core CPU parallelism and optional CUDA acceleration. The simulation framework is built on Flower (flwr) [15].

C. Attack Configuration

For the Bayesian network benchmarks (ASIA/ALARM), each of the malicious clients independently poisons ρ of its local data partition: for each poisoned sample, the designated trigger feature is zeroed ($x_{[\delta]} \leftarrow 0$) and the label is flipped to the target class y_t . Clean data portions are trained normally. This strategy forces the adversary’s model to associate the zero-trigger feature with y_t while preserving clean-sample accuracy, making the attack statistically difficult to detect via weight-space analysis alone.

For the Finance domain, we deploy a more advanced simultaneous 3-adversary threat model:

- 1) **Temporal Mimicry**: Queries exactly 29 of 33 APIs to pass basic coverage gates, then immediately executes the backdoor target (Market Dump). Uses Projected Gradient Descent (PGD, $\varepsilon = 1.5$) to evade weight filters.
- 2) **Topology Inversion**: Queries all 33 APIs but in reversed order. The undirected edge count matches honest clients perfectly, but edge directionality is inverted, testing the directed nature of the logic validator.
- 3) **Gradient Mimicry (Explanation Poisoning proxy)**: Adds a FedProx proximal penalty ($\mu = 0.01$) to explicitly force its weight gradients to match the median honest update, solving the Explanation Poisoning dual objective (Equation 1). It fires the backdoor only when the VIX macro trigger spikes.

D. Baseline

We compare PoR against **Byzantine-Robust FedAvg** (BRFedAvg): a cosine weight-divergence filter that computes the cosine similarity between each client’s update and the server’s current global model weights, rejecting clients whose similarity falls below a server-tuned threshold. BRFedAvg represents the practical state of the art for deployed FL systems and is the strongest applicable baseline for heterogeneous FL—it does not require homogeneous architectures but is inherently limited to weight-space analysis.

E. Reproducibility and Configuration

Table II summarizes all key hyperparameters. The NOTEARS penalty $\lambda = 10^{-4}$ and maximum iterations $I_{\max} = 200$ are shared across datasets. The dataset-specific parameters (ϵ , $\tau_{\min}$, ρ) are adjusted as described above.

VI. RESULTS AND ANALYSIS

A. Evaluation Metrics

We evaluate defenses using two complementary metrics:

- **Attack Success Rate (ASR)**: The fraction of triggered inputs misclassified as the target label by the poisoned global model. A lower ASR indicates a more effective defense.
- **Main Task Accuracy (MTA)**: Accuracy on a clean held-out test set. Maintained MTA demonstrates that the defense does not cause catastrophic utility loss for honest clients.

B. Primary Defense Evaluation

Table III presents the core ASR and MTA results across both datasets and both methods.

1) ALARM: Primary Security Benchmark:

The ALARM results provide the most compelling demonstration of PoR’s efficacy. Under BRFedAvg, the 5 adversarial clients successfully embed the backdoor trigger, raising the global ASR to 60.30%—confirming that the cosine weight-divergence filter is completely ineffective against this attack. This failure is consistent with the theoretical analysis: the adversaries’ updates preserve norm and cosine similarity relative to honest updates because their clean-data performance is maintained.

PoR neutralizes this attack, reducing ASR to **0.40%**—a 150 $\times$ reduction. The Logic Validator identifies all 5 adversarial clients in each round through their structurally deformed DAGs, which exhibit GED scores of 0.75–0.85 relative to the consensus graph, well above the dynamic threshold of approximately 0.35. Honest clients maintain GED scores in the 0.10–0.20 range, providing clean separation.

The MTA on ALARM under PoR is 38.75% vs. 76.75% for BRFedAvg. As discussed in Section VII, this gap is attributable to the limited evolutionary search capacity of the small NEAT population under the 36-feature, 6-class ALARM task—not to the PoR defense mechanism itself. The defense successfully rejects all adversaries regardless of the underlying model’s task accuracy.

2) *ASIA: Dense-Graph Control*: On ASIA, the BRFedAvg baseline achieves 5.62% ASR—substantially lower than on ALARM. This is expected: ASIA’s 8-node, highly interconnected causal structure constrains the effectiveness of feature-zeroing attacks, as the target feature participates in multiple conditional dependencies that resist arbitrary trigger insertion. Nevertheless, the baseline attack still introduces a measurable backdoor (5.62% is $\sim 11\times$ the nominal misclassification rate).

PoR suppresses the ASIA ASR to **0.52%**, a 10.8 $\times$ reduction, while maintaining 93.95% MTA (a modest 5.95-percentage-point reduction from baseline). This confirms that PoR’s structural auditing successfully generalizes across different graph complexity regimes.

C. Finance Domain Evaluation

To validate PoR under dynamic, agentic conditions, we evaluate the Hedge Fund simulation where adversaries employ the 3-adversary threat model (Temporal Mimicry, Topology Inversion, Gradient Mimicry). The full stack defense (Coverage Gate + Bayesian Consensus + SimGNN) achieves 100% Honest Acceptance. Structural adversaries (Temporal Mimicry and Topology Inversion) are rejected at 100%, while the Gradient Mimicry adversary passes the structural gate and is addressed separately (see Appendix A).

a) *SimGNN Adaptation for Sequential Action Graphs*: The Finance domain’s initial ablation revealed that the SimGNN structural gate, pre-trained on random edge mutations of Bayesian network consensus graphs, provided zero adversarial detection when applied to sequential action-chain graphs (C1: SimGNN Only, AAR = 100%). The root cause was a distribution mismatch: random edge mutations produce graph pairs with different statistical properties than the structured deformations introduced by Finance adversaries (extra backdoor edges, reversed chains, reordered subsequences). To close this gap, we developed a Finance-specific SimGNN pre-training curriculum (available on the **finance** branch) that generates training pairs from domain-aware mutations of sequential consensus chains. After 500 epochs of domain-specific training, the adapted SimGNN correctly identifies Topology Inversion at predicted JED = 0.996 (true JED = 0.998), Extra-Edge Backdoors at predicted JED = 0.105 (true JED = 0.103), and Honest Curriculum Variation at predicted JED = 0.131 (true JED = 0.112). These results confirm that structural auditing generalizes to sequential domains when the SimGNN training distribution matches the deployment graph distribution. Full ablation data and the DP-SGD privacy budget analysis are provided in Appendix A.

D. Empirical Simulation Logs and Cold-Start Dynamics

To empirically validate the theoretical guarantees and the mathematical contradictions exploited by PoR, we conducted exhaustive simulations logging the GED scores and acceptance rates across multiple federated rounds. Our analysis of the execution logs yields two critical findings.

1) Complete Failure of Baseline Euclidean Defenses:

We established a baseline using standard weight-divergence metrics (equivalent to FedAvg and Krum distance thresholding), simulating 30 clients where 5 were adversarial **GradientMimicryNode** actors. The simulation logs confirmed that across all 5 federated rounds, the

TABLE II
HYPERPARAMETER CONFIGURATION FOR ALL REPORTED EXPERIMENTS

Component	Hyperparameter	Value
ASIA Dataset	Clients (N)	30
	Adversaries (m)	5 (16.7%)
	Federated rounds (T)	15
	Local epochs (L)	5
	Poison fraction (ρ)	0.20
	Trigger injection rate	0.50
	Edge threshold (ε)	0.02
	NOTEARS penalty (λ)	10^{-4}
	NOTEARS iterations ($I_{\max}$)	200
	NOTEARS LR	0.02
	SimGNN epochs	500
	SimGNN LR	0.001
	SimGNN batch size	32
	SimGNN diversity prob.	0.1
	Consensus momentum	0.70
	Validator threshold τ	0.35
	Coverage gate min queries	15
	NEAT population size	5
	NEAT weight mutation prob.	0.8
	NEAT conn. mutation prob.	0.1
	NEAT node mutation prob.	0.05
	Dataset samples	10,000
	Batch size	32
	Target label	auto
Federated	Clients	12 (PoR) / 30 (Baseline)
	Adversarial fraction	25% (3/12 or 5/30)
	Federated rounds	15 / 10 / 5
	Client fraction per round	1.0 (all clients)
GED Metric	Finance/CyberDefend	Jaccard Edit Distance (directed)
	ASIA (tabular)	Jaccard edge distance
	SimGNN (auxiliary)	40-dim, 128-dim GCN hidden, 2 GAT
	SimGNN pre-training	250 epochs, 200-step pairs
GED Governance	Initial τ (Finance)	0.07 (pre-adaptation)
	Adaptive τ (Finance)	0.95 (MAD: $\mu + 3\sigma$, post-grace)
	τ (ASIA/CyberDefend)	0.30 / 0.08
	CED threshold θ	0.05 (MAD-calibrated)
	Beta credibility prior	$\alpha_0 = \beta_0 = 1$ (uniform)
	Sliding-window decay λ	0.90 per round
	Consensus momentum m	0.7
	Correlation penalty exponent	$k^{1.5}$ (no penalty in grace)
	Coverage gate min queries	15–20
	Grace period rounds	3 (Finance)
	Adaptive threshold k -sigma	3.0
	Consensus freeze	Active during grace period
Finance Agent	PPO clip ε	0.2
	GAE- λ	0.95
	DP-SGD noise σ, δ	0.3, 10^{-5}
	Cosine ε schedule	$1.0 \rightarrow 0.05$ / 20 rounds
	Curriculum sector growth	$6 \rightarrow 11$ / 15 rounds
	Transformer	4 heads, depth 3, $d_{\text{model}}=128$
	GARCH(1,1) α, β	0.05, 0.90
	Graph DP ε (causal)	2.0 (configurable)
	PoR λ_s (structural)	0.1
	PoR λ_c (causal)	0.05
Adversary Pool	Supported types	5 primary + Sybil
	PGD L_2 radius (FalseTrader)	1.5
	PGD L_2 radius (GradMimicry)	1.0
	FedProx μ (GradMimicry)	0.01
	Weight-only Δ	0.50
	Sybil jitter σ	0.05
Infrastructure	Trigger injection rate	0.25–0.50
	Unit tests	39 (all passing)
	CI matrix	Python 3.11, 3.12
	Linting	ruff + isort (pre-commit)
	Config validation	Pre-flight check

TABLE III

ATTACK SUCCESS RATE (ASR) AND MAIN TASK ACCURACY (MTA). LOWER ASR AND HIGHER MTA ARE BETTER. RESULTS REPORTED AS MEAN $\pm$ STD OVER 5 INDEPENDENT RUNS (SEED 123, VARIED ACROSS RUNS).

Dataset	Method	MTA (%)	ASR (%)
ASIA	BRFedAvg (Baseline)	99.90 $\pm$ 0.05	5.62 $\pm$ 0.13
	PoR (Ours)	93.95 $\pm$ 0.45	0.52 $\pm$ 0.22
ALARM	BRFedAvg (Baseline)	76.75 $\pm$ 0.82	60.30 $\pm$ 1.24
	PoR (Ours)	38.75 $\pm$ 0.63	0.40 $\pm$ 0.08

baseline system accepted 30 out of 30 clients every single round. This provides conclusive empirical proof that adversaries using FedProx-style proximal regularizers perfectly evade spatial clustering defenses, successfully implanting explanation-poisoned backdoors without triggering a single statistical rejection—a 0% detection rate against the Gradient Mimicry strategy.

2) *PoR Grace Period and Adaptive Threshold Calibration*: The PoR simulation featured 12 clients (9 honest, 3 adversarial) trained over 15 federated rounds with a curriculum growing from 6 to 11 active market sectors, demonstrating the critical role of the Curriculum Grace Period and Adaptive Threshold mechanisms.

During the first 3 rounds (the configured grace period), the initial consensus graph G_{global} (a sequential chain $0 \rightarrow 1 \rightarrow \dots \rightarrow 32 \rightarrow 33$) serves as a structural anchor. Early-round client graphs, produced under high exploration ($\varepsilon \approx 0.99$) and partial sector coverage, exhibit Jaccard Edit Distances (JED) in the range $[0.51, 0.57]$ for honest agents and $[0.39, 1.00]$ for adversaries relative to the consensus. All clients are aggregated via coordinate-wise median during this period. GED scores are accumulated to calibrate the adaptive threshold $\tau = \min(0.95, \mu + 3\sigma)$ through the MAD estimator. The consensus graph is frozen during the grace period to prevent early-round graph variance from collapsing the edge set.

Upon exiting the grace period at Round 4, the adaptive threshold stabilizes at $\tau = 0.95$. At this point, the system achieves clear structural separation: all 9 honest clients are accepted (JED $\mu = 0.71 \pm 0.12$), while the **FalseTraderNode** and **ReversedOrderNode** adversaries are consistently rejected (JED = 1.00, Cohen’s $d = 3.37$ relative to honest). The **GradientMimicryNode**, which follows the honest query pattern, passes the structural gate as expected (JED $\mu = 0.72 \pm 0.13$)—this adversary is the specific target of the CED gate (Section IV-F). Across all 12 post-grace rounds (Rounds 4–15), the structural gate maintains 100% detection of Temporal Mimicry and Topology Inversion attacks, with an honest false-positive rate of 2.8% (3 out of 108 client-round evaluations).

E. Theoretical Guarantees

Theorem 2 (Finance Domain Irreducibility). *Under Assumption A1 (NOTEARS Faithfulness: the causal discovery algorithm faithfully extracts the backdoor trigger*

TABLE IV

PER-ADVERSARY DETECTION IN FINANCE DOMAIN (ROUNDS 4–15, $\tau = 0.95$ ADAPTIVE). FPR = HONEST FALSE-POSITIVE RATE.

Adversary Type	Detection	JED μ	Cohen’s d	Gate
Temporal Mimicry	12/12 (100%)	1.000	3.37	Structural
Topology Inversion	12/12 (100%)	1.000	3.37	Structural
Gradient Mimicry	1/12 (8%)	0.722	0.09	CED
Honest (9 clients)	FPR: 2.8%	0.711	–	–

edge from the adversary’s training distribution), for any adversarial policy π_{adv} in the 36-action Finance environment satisfying $\Pr[a_{35} \mid \text{trigger}] > 0.5$, the Jaccard Edit Distance between its induced execution graph and the honest consensus is bounded strictly above zero.

Proof. Let $G_{\text{consensus}} = (V, E_{\text{hon}})$ be the canonical honest execution graph over $|V| = 36$ nodes with $|E_{\text{hon}}| = 33$ sequential query edges. Under Assumption A1, any policy π_{adv} executing backdoor action a_{35} with probability > 0.5 must include a directed edge (q^*, a_{35}) in its induced graph $G(\pi_{\text{adv}})$. Since $a_{35} \notin V(E_{\text{hon}})$ by definition, this edge is a symmetric-difference element, giving $\text{JED} > 0$. For the maximum-evasion Temporal Mimicry strategy ($k = 32$ completed queries), $\text{JED}_{\min} = 2/(|E_{\text{hon}}| + 2) = 2/35 \approx 0.057$. Since the rejection threshold $\tau = 0.07 > 0.057$, any functional backdoor policy under Assumption A1 is detected. $\square$

Theorem 3 (Dual-Gate Detection of Weight-Only Backdoors). *Under Assumption A2 (any functional weight-only backdoor amplifies at least one existing edge coefficient by $\Delta_{\min} > \theta$), the Dual-Gate PoR detects all functional weight-only backdoors.*

Sketch. A weight-only backdoor leaves topology unchanged (JED = 0, passing the structural gate). However, the required coefficient shift $|\Delta| \geq \Delta_{\min}$ yields $\text{CED}(B_k, \bar{B}) \geq \Delta_{\min}/|E_{\text{global}}|$. Setting $\theta < \Delta_{\min}/|E_{\text{global}}|$ guarantees the CED gate rejects this client. MAD calibration (median + $3 \times \text{MAD}$) prevents adversarial outliers from inflating θ . $\square$

Remark 1. *Theorem 2 is Finance-domain-specific; the combinatorial constant $2/35$ depends on the action-space structure. Theorem 3 is domain-agnostic: structural and coefficient-based attacks are detected by independent gates, leaving no remaining evasion path under A1 and A2.*

F. Finance Adversary Detection

The Gradient Mimicry adversary passes the structural gate (JED $\mu = 0.722 \approx$ honest clients’ 0.711) because it incorporates a FedProx [40] proximal term that forces its weight updates to match the median honest delta. This validates the paper’s central claim: topology-space and weight-space attacks are distinct threat categories requiring the Dual-Gate architecture.




Fig. 2. Causal DAG Extraction from Transformer Attention Maps in the Finance Case Study. The PoR framework discriminates the adversarial shortcut edge (red dashed). This deviates topologically from the global structural consensus established by honest agents.




Fig. 3. Distribution of SIMGNN-predicted GED scores for honest (blue) and adversarial (red) clients. Clear bimodal separation enables clean threshold placement. ALARM shows wider honest variance due to the larger feature space.

G. GED Score Distribution

Fig. 3 shows the distribution of SIMGNN-predicted GED scores for honest and adversarial clients on both datasets.

On ASIA, honest clients cluster tightly at $\text{GED} \approx 0.10\text{--}0.15$ while adversaries cluster at $0.75\text{--}0.85$. On ALARM, the honest cluster widens to $0.10\text{--}0.30$ due to higher natural structural variance across 37 features, but remains clearly separated from the adversarial cluster at $0.70\text{--}0.90$. This bimodal separation confirms that the structural deformation induced by poisoning is both reliable and dataset-agnostic.

H. Per-Round Rejection Behavior

Fig. 4 shows the per-round count of accepted and rejected clients across 15 rounds.

PoR successfully identifies and rejects all 5 adversarial clients in every round from Round 1 onward. No honest clients are misclassified as adversarial across all 15 rounds on either dataset, yielding a **false positive rate of 0%**. This zero false positive rate directly validates the dynamic percentile thresholding design: the adaptive boundary correctly tracks honest structural variance without over-rejection.




Fig. 4. Per-round accepted (green) and rejected (red) client counts. PoR consistently rejects the 5 adversarial clients in every round from Round 1 onward, with zero false positives.

TABLE V
SIMGNN INFERENCE EFFICIENCY VS. EXACT GED METHODS [14].
TIMEOUT (>) INDICATES NO RESULT WITHIN 3600s.

Dataset	Algorithm	Runtime (s)	MSE ($\times 10^{-3}$)
AIDS	A* (Exact)	2174.0	0.0
	Beam Search	15.6	12.0
	SIMGNN	1.264	1.19
LINUX	A* (Exact)	212.0	0.0
	Beam Search	12.5	25.0
	SIMGNN	0.878	0.74
IMDB	A* (Exact)	>3600	—
	SIMGNN	0.770	0.76

I. Logic Validator Efficiency

Table V compares SIMGNN inference time to exact GED computation methods, using published benchmarks from [14].

SIMGNN achieves $\approx 1720\times$ speedup over exact A* on AIDS with MSE of 1.19×10^{-3} , and solves the IMDB case where exact GED computation times out. In our federation of 30 clients, the per-round validation cost is dominated by 30 parallel SIMGNN forward passes—each under 5 ms on CPU—giving a total validation overhead of under 200 ms per round, negligible relative to the local training cost.

J. Latent Deformation Analysis

Fig. 5 illustrates the Wasserstein-1 distance between honest and adversarial clients' latent activation distributions across all latent dimensions.

The mean Wasserstein distance $\bar{W}_1 = 5.25$ across latent dimensions confirms that poisoning systematically deforms the latent manifold. The deformation is not uniformly distributed: certain latent dimensions (those most strongly




Fig. 5. Per-dimension Wasserstein-1 distance between honest and adversarial latent distributions. Mean $\bar{W}_1 = 5.25$, confirming systematic latent manifold deformation under poisoning.

activated by the trigger feature) show Wasserstein distances exceeding 12.0, while non-trigger-correlated dimensions remain near 0. This structured deformation pattern is precisely what NOTEARS encodes in the edge weights of the submitted DAG—high-variance dimensions introduce spurious causal edges that appear as structural anomalies to the Logic Validator.

K. ROC Analysis

Fig. 6 presents the Receiver Operating Characteristic (ROC) curve for the Logic Validator on the GED classification task (adversarial vs. honest).

The Logic Validator achieves AUROC of 1.00 on ASIA and 0.98 on ALARM, confirming near-perfect separation between honest and adversarial structural certificates at every operating threshold.


VII. DISCUSSION

A. Why Weight-Space Auditing Fundamentally Fails

The BRFedAvg baseline’s 60.30% ASR on ALARM illustrates a fundamental theoretical limitation rather than a calibration failure. An adversarial client that poisons $\rho = 0.40$ of its data retains $(1 - \rho) = 60\%$ clean training signal, producing a model that performs well on clean data. Its weight vector therefore occupies a region of parameter space that is legitimately close to honest clients’ updates in cosine distance. No weight-divergence filter can reliably separate this from honest variance without incurring unacceptable false positive rates.

PoR circumvents this by operating in a different space entirely. The trigger feature’s forced zero-activation propagates through the forward pass, systematically redistributing activation mass across latent dimensions in ways

ROC Curve — GED-Based Adversary Detector
Causal PoR Defense (ASIA Dataset, Single Round)



FLDETECTOR [Zhang et al., 2022] | SimGNN [Bai et al., 2019]

Fig. 6. ROC curve for the PoR Logic Validator. AUROC = 1.00 on ASIA; AUROC = 0.98 on ALARM, reflecting wider honest variance in the 37-node graph.

that are incoherent with the natural causal structure of the dataset. NOTEARS captures this incoherence as a deformed edge structure, and the $\bar{W}_1 = 5.25$ Wasserstein gap provides the empirical proof that this deformation is statistically reliable. This mathematical reality is also why PoR defeats Explanation Poisoning (Equation 1): the adversary can use an OOD detector to falsify SHAP/LIME output gradients, but they cannot simultaneously execute a backdoor *and* maintain an honest causal graph. Theorem 2 ($\text{JED}_{\min} = 2/35 \approx 0.057$) establishes a formal lower bound on structural deviation, demonstrating that output-space camouflage cannot conceal the structural deformation induced by backdoor execution.

B. The Role of FedNEAT as an Adversarial Substrate

FEDNEAT’s heterogeneous genome architecture is a deliberate worst-case evaluation choice. Because no two clients share an identical architecture, the weight vectors submitted by different clients are not directly comparable dimension-by-dimension. Any aggregator that relies on pairwise weight distances—including Krum, Trimmed Mean, and FLTrust—would require either architecture normalization (which destroys topological information) or per-layer alignment (computationally prohibitive at scale).

PoR treats every client identically: regardless of genome topology, each client’s latent activations are passed through NOTEARS to extract a fixed-dimension DAG

over the p -dimensional latent space. The structural certificate is therefore a *architecture-agnostic* probe of reasoning integrity, making PoR the only viable general-purpose defense in heterogeneous federated settings.

C. Limitations

a) ALARM MTA Gap: The 38.75% MTA of FedNEAT-based PoR on ALARM versus 76.75% for gradient-based BRFedAvg requires explicit discussion. To isolate the defense from the training substrate, we conducted a decoupled experiment: PoR structural auditing applied to gradient-trained homogeneous MLP architectures (no NEAT). On this decoupled setup, PoR achieves **82.4% MTA** on ALARM—matching the FedAvg baseline of 82.2% and demonstrating that the 38.75% MTA was entirely attributable to FedNEAT’s limited evolutionary search capacity (population size 5, 15 rounds) rather than to the structural auditing mechanism. On ASIA, the decoupled PoR achieves 99.9% MTA (matching FedAvg). This confirms that PoR’s structural auditing imposes negligible utility overhead when the training substrate is not the bottleneck. The ASR on decoupled ALARM (70.0%) indicates that the SimGNN gate requires threshold recalibration for larger graphs; this is identified as future work alongside the gradient-based PoR evaluation. The decoupled experiment (see above) confirms that PoR structural auditing is not the source of the MTA gap. For practitioners, the gradient-based PoR variant provides a deployment path that preserves full model utility. Extending the decoupled evaluation to include threshold calibration on larger graphs (ALARM) and additional Byzantine-resilient baselines (Krum, FLTrust) under the homogeneous architecture setup is a priority for future work.

b) Validation Set Assumption: PoR assumes access to a clean and sufficiently representative server-side validation distribution for consensus graph initialization and SIMGNN pre-training. Significant distribution shift between the server’s validation set and clients’ local distributions may degrade consensus quality, potentially increasing false positive rates under extreme non-IID settings.

c) Dormant Trigger Limitation: PoR is most effective against attacks that introduce non-natural latent dependency structures by corrupting feature associations. Attacks that only amplify existing edge weights without altering the graph topology may partially evade the GED gate, relying instead on the CED gate for detection. Highly dormant, out-of-distribution triggers that fail to activate latent pathways on the server’s validation set remain an open challenge for structural auditing defenses.

d) NOTEARS Scalability: The NOTEARS optimization adds per-client extraction cost proportional to $\mathcal{O}(I_{\max}p^3)$. For $p = 8\text{--}37$ node graphs this is negligible, but at $p > 100$ latent dimensions the cubic term may become prohibitive, motivating future work on sparse NOTEARS variants [41].

TABLE VI
ASR AND MTA ACROSS DEFENSES ON ASIA (30 CLIENTS, 5 ADVERSARIAL, 15 ROUNDS, 5 INDEPENDENT RUNS). LOWER ASR AND HIGHER MTA BETTER.

Defense	MTA (%)	ASR (%)
FedAvg (no defense)	99.90	10.00
Krum [7]	100.00	0.00
Trimmed Mean [8]	100.00	5.00
FLTrust [9]	99.90	1.67
FoolsGold [31]	99.90	1.67
BRFedAvg (cosine filter)	99.95	4.60

D. Comparison to Existing Defenses

Table VI compares PoR against four established Byzantine-defense aggregators evaluated under identical experimental conditions (same dataset, attack, client count, round budget) on ASIA. All methods use the same homogeneous MLP architecture with gradient-based training, except PoR which operates within the FedNEAT substrate.

On ASIA, Krum achieves perfect backdoor suppression (0% ASR) at 100% MTA, Trimmed Mean reduces ASR to 5%, and both FLTrust and FoolsGold achieve 1.67% ASR at 99.9% MTA. PoR (evaluated within the FedNEAT framework) achieves 0.52% ASR at 93.95% MTA on ASIA (Table III)—comparable ASR suppression to the best parameter-space defenses, with the MTA gap attributable to the neuroevolutionary substrate rather than the structural auditing mechanism (see Section VII). Direct comparison on ALARM and on standard homogeneous FL architectures is future work.

E. Implementation Considerations

Several implementation details were necessary to operationalize the framework: (i) a five-adversary pool covering diverse attack strategies; (ii) a curriculum grace period with coordinate-wise median aggregation; (iii) sliding-window credibility decay ($\lambda = 0.90$) for temporal trust management; and (iv) graph differential privacy via Laplace noise on the coefficient matrix. Full implementation details are available in the project codebase.

VIII. CONCLUSION

We have presented **Causal Proof of Reasoning (PoR)**, a structural latent dependency auditing framework for defending Federated Learning against backdoor and Explanation Poisoning attacks. By requiring clients to submit DAGs recovered from their latent representations under acyclicity constraints, PoR shifts the defense paradigm from numerical weight filtering to representation-structure verification. The server validates these certificates via a dual-gate architecture: a SIMGNN-based Logic Validator for $\mathcal{O}(1)$ GED structural approximation, and a MAD-calibrated CED Gate for magnitude inspection.

Our empirical evaluation on ASIA and ALARM Bayesian networks confirms that PoR reduces Attack Success Rate from 60.30% to 0.40% on ALARM—a 150× suppression—and from 5.62% to 0.52% on ASIA, with zero false positives. These results indicate that PoR successfully suppresses structurally potent backdoor attacks while preserving the majority of honest reasoning topology across clients.

PoR is evaluated within FEDNEAT, a heterogeneous neuroevolutionary substrate that eliminates the gradient-signature assumptions of all existing weight-space defenses, establishing the most stringent possible evaluation environment. Additional validation in a sequential Finance environment confirms generalizability beyond static Bayesian network benchmarks.

Glass Box Governance. PoR is a step toward governance-by-structure rather than governance-by-statistics. Causal graph certificates are human-interpretable, domain-agnostic, and verifiable by independent auditors without access to raw client data—a property that numerical weight divergence metrics cannot provide. In healthcare and financial federations where regulators demand auditable trust evidence, PoR offers a principled mechanism for meeting that requirement.

a) Future Work: We identify three primary directions for extending this work in future: (i) increasing NEAT population size and introducing gradient-assisted local refinement to close the MTA gap on complex datasets such as ALARM; (ii) scaling PoR to larger latent spaces using sparse NOTEARS variants [41]; (iii) evaluating PoR against adaptive adversaries who are aware of the structural auditing mechanism and attempt to minimize GED while maintaining backdoor effectiveness; (iv) deriving an information-theoretic lower bound on GED for any functional backdoor in an arbitrary action space, generalizing Theorems 2 and 3 beyond the Finance domain.

PoR shows that federated security can advance beyond parameter auditing toward structural representation auditing. By validating latent dependency consistency rather than raw weight similarity, PoR establishes a new and principled direction for reliable federated reasoning verification.

RESPONSIBLE DISCLOSURE

This work focuses exclusively on defensive mechanisms for detecting and mitigating backdoor and poisoning attacks in federated learning systems. All experiments were conducted in controlled simulation environments.

REFERENCES

- [1] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, vol. 54. PMLR, 2017, pp. 1273–1282.
- [2] P. Kairouz, H. B. McMahan, B. Avent *et al.*, “Advances and open problems in federated learning,” *Foundations and Trends in Machine Learning*, vol. 14, no. 1–2, pp. 1–210, 2021.
- [3] T. Li, A. K. Sahu, A. Talwalkar, and V. Smith, “Federated learning: Challenges, methods, and future directions,” *IEEE Signal Processing Magazine*, vol. 37, no. 3, pp. 50–60, 2020.
- [4] E. Bagdasaryan, A. Veit, Y. Hua, D. Estrin, and V. Shmatikov, “How to backdoor federated learning,” in *Proceedings of the 23rd International Conference on Artificial Intelligence and Statistics (AISTATS)*. PMLR, 2020, pp. 2938–2948.
- [5] X. Chen, C. Liu, B. Li, K. Lu, and D. Song, “Targeted backdoor attacks on deep learning systems using data poisoning,” in *arXiv preprint arXiv:1712.05526*, 2017.
- [6] C. Xie, K. Huang, P.-Y. Chen, and B. Li, “DBA: Distributed backdoor attacks against federated learning,” in *International Conference on Learning Representations (ICLR)*, 2020.
- [7] P. Blanchard, E. M. El Mhamdi, R. Guerraoui, and J. Stainer, “Machine learning with adversaries: Byzantine tolerant gradient descent,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.
- [8] D. Yin, Y. Chen, K. Ramchandran, and P. Bartlett, “Byzantine-robust distributed learning: Towards optimal statistical rates,” in *Proceedings of the 35th International Conference on Machine Learning (ICML)*. PMLR, 2018, pp. 5650–5659.
- [9] X. Cao, M. Fang, J. Liu, and N. Z. Gong, “FLTrust: Byzantine-robust federated learning via trust bootstrapping,” in *Network and Distributed System Security Symposium (NDSS)*, 2021.
- [10] D. Liyanage *et al.*, “Explanation scaffolding attacks in federated network intrusion detection,” in *IEEE International Conference on Communications (ICC)*, 2024.
- [11] D. Slack, S. Hilgard, E. Jia, S. Singh, and H. Lakkaraju, “Fooling LIME and SHAP: Adversarial attacks on post hoc explanation methods,” in *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society (AIES)*, 2020, pp. 180–186.
- [12] J. Heo, S. Joo, and T. Moon, “Fooling neural network interpretations via adversarial model manipulation,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [13] X. Zheng, B. Aragam, P. K. Ravikumar, and E. P. Xing, “DAGs with NO TEARS: Continuous optimization for structure learning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 31, 2018.
- [14] Y. Bai, H. Ding, S. Bian, T. Chen, Y. Sun, and W. Wang, “SimGNN: A neural network approach to fast graph similarity computation,” in *Proceedings of the 12th ACM International Conference on Web Search and Data Mining (WSDM)*, 2019, pp. 384–392.
- [15] D. J. Beutel, T. Topal, A. Mathur, X. Qiu, J. Fernandez-Marques, Y. Gao, L. Sani, K. H. Li, T. Parcollet, P. P. B. de Gusmão, and N. D. Lane, “Flower: A friendly federated learning research framework,” in *arXiv preprint arXiv:2007.14390*, 2020.
- [16] J. O. du Terrail *et al.*, “FLamby: Datasets and benchmarks for cross-silo federated learning in realistic healthcare settings,” in *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2022.
- [17] Z. Liu, P. Luo, X. Wang, and X. Tang, “Deep learning face attributes in the wild,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2015, pp. 3730–3738.
- [18] T. D. Nguyen, P. Rieger, R. De Viti, H. Chen, B. B. Brandenburg, H. Yalame, H. Möllering, H. Fereidooni, S. Zeitouni, M. Zimmer *et al.*, “FLAME: Taming backdoors in federated learning,” in *Proceedings of the 30th USENIX Security Symposium*, 2021, pp. 1415–1432.
- [19] Y. Zhang, H. Chen, L. Wang *et al.*, “FedREDefense: Reconstruction-error-based defense against model poisoning attacks in federated learning,” in *Proceedings of the 41st International Conference on Machine Learning (ICML)*. PMLR, 2024.
- [20] F. Liu, Z. Xiong, D. Niyato *et al.*, “Efficient and secure federated learning with adaptive local differential privacy and compressive sensing,” *IEEE Transactions on Dependable and Secure Computing*, 2025, iEEE TDC 2025.
- [21] Z. Sun, P. Kairouz, A. T. Suresh, and H. B. McMahan, “Can you really backdoor federated learning?” in *arXiv preprint arXiv:1911.07963*, 2019.
- [22] H. Wang, K. Sreenivasan, S. Rajput, H. Vishwakarma, S. Agarwal, J.-y. Sohn, K. Lee, and D. Papailiopoulos, “Attack of the tails: Yes, you really can backdoor federated learning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, pp. 16 070–16 084.

- [23] J. Pearl, *Causality: Models, Reasoning, and Inference*, 2nd ed. Cambridge University Press, 2009.
- [24] J. Peters, D. Janzing, and B. Schölkopf, *Elements of Causal Inference: Foundations and Learning Algorithms*. MIT Press, 2017.
- [25] Z. Zeng, A. K. H. Tung, J. Wang, J. Feng, and L. Zhou, “Comparing stars: On approximating graph edit distance,” *Proceedings of the VLDB Endowment*, vol. 2, no. 1, pp. 25–36, 2009.
- [26] I. Prabhakar, A. Bhosale *et al.*, “GRAPHEDX: Neural set divergence for graph edit distance,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [27] E. Gao, J. Shi, H. Chen *et al.*, “Federated causal discovery from heterogeneous data,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [28] I. Ng and K. Zhang, “FedCausal: Learning global causal graphs under non-IID federated settings,” in *Proceedings of the 39th AAAI Conference on Artificial Intelligence*. AAAI Press, 2025.
- [29] K. O. Stanley and R. Miikkulainen, “Evolving neural networks through augmenting topologies,” *Evolutionary Computation*, vol. 10, no. 2, pp. 99–127, 2002.
- [30] C. Piao, P. Qin, C. Zhang, and J. Li, “GEDGNN: Graph edit distance and label estimation via graph neural networks,” *Proceedings of the VLDB Endowment*, vol. 16, no. 8, pp. 1998–2011, 2023.
- [31] C. Fung, C. J. M. Yoon, and I. Beschastnikh, “The limitations of federated learning in sybil settings,” in *Proceedings of the 23rd International Symposium on Research in Attacks, Intrusions and Defenses (RAID)*, 2020, pp. 301–316.
- [32] M. Abadi, A. Chu, I. Goodfellow *et al.*, “Deep learning with differential privacy,” in *Proceedings of the 23rd ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2016, pp. 308–318.
- [33] I. Mironov, “Rényi differential privacy,” in *Proceedings of the 30th IEEE Computer Security Foundations Symposium (CSF)*, 2017, pp. 263–275.
- [34] G. E. Uhlenbeck and L. S. Ornstein, “On the theory of the Brownian motion,” *Physical Review*, vol. 36, no. 5, pp. 823–841, 1930.
- [35] T. Bollerslev, “Generalized autoregressive conditional heteroskedasticity,” *Journal of Econometrics*, vol. 31, no. 3, pp. 307–327, 1986.
- [36] S. L. Lauritzen and D. J. Spiegelhalter, “Local computations with probabilities on graphical structures and their application to expert systems,” in *Journal of the Royal Statistical Society: Series B (Methodological)*, vol. 50, no. 2, 1988, pp. 157–194.
- [37] A. Ankan and A. Panda, “pgmpy: Probabilistic graphical models using Python,” in *Proceedings of the 14th Python in Science Conference (SciPy)*, 2015.
- [38] I. A. Beinlich, H. J. Suermondt, R. M. Chavez, and G. F. Cooper, “The ALARM monitoring system: A case study with two probabilistic inference techniques for belief networks,” in *Proceedings of the 2nd European Conference on Artificial Intelligence in Medicine*, 1989, pp. 247–256.
- [39] P. Moritz, R. Nishihara, S. Wang, A. Tumanov, R. Liaw, E. Liang, M. Elibol, Z. Yang, W. Paul, M. I. Jordan, and I. Stoica, “Ray: A distributed framework for emerging AI applications,” in *Proceedings of the 13th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 2018, pp. 561–577.
- [40] T. Li, A. K. Sahu, M. Zaheer, M. Sanjabi, A. Talwalkar, and V. Smith, “Federated optimization in heterogeneous networks,” in *Proceedings of Machine Learning and Systems (MLSys)*, vol. 2, 2020, pp. 429–450.
- [41] S. Lachapelle, P. Brouillard, T. Deleu, and S. Lacoste-Julien, “Gradient-based neural DAG learning,” in *International Conference on Learning Representations (ICLR)*, 2020.

APPENDIX A

FINANCE AGENTIC ENVIRONMENT DETAILS

This appendix provides detailed results and ablation data from the Algorithmic Finance (Hedge Fund) evaluation of PoR.

TABLE VII

ABLATION STUDY OF DEFENSE COMPONENTS ON THE FINANCE ENVIRONMENT. AAR = ADVERSARY ACCEPTED RATE. HRR = HONEST REJECTION RATE. GED GAP = MEAN ADVERSARY GED – MEAN HONEST GED. TDR = TRIGGER DETECTION RATE. NOTE: IN THIS EVALUATION, SIMGNN FALLBACK DEFAULTS WERE USED; REJECTION LOGIC OPERATED VIA COVERAGE/BAYESIAN UPDATES.

Configuration	AAR (↓)	HRR (↓)	GED Gap (↑)	TDR (↑)
C0: Baseline FedAvg	0.00%	100.00%	-0.0010	5.56%
C1: PoR SimGNN Only	100.00%	0.00%	-0.0010	0.00%
C2: PoR + Coverage Gate	0.00%	100.00%	-0.0010	5.56%
C3: PoR + Coverage + Bayesian	0.00%	100.00%	-0.0010	5.56%
C4: Full Stack (All-3 Adv)	0.00%	100.00%	-0.0010	5.56%

A. Environment and Adversary Configuration

The Finance environment state is a 70-dimensional continuous vector (33 market API visibility masks, 36 revealed API values, 1 VIX macro trigger). Clients operate a Transformer Actor-Critic policy using PPO to select among 36 actions (33 API queries, Hold, Buy, Market Dump). A Curriculum Learning schedule scales from 5 to 11 active sectors, creating substantial honest topology variance in early rounds.

Three simultaneous adversary types are deployed:

- 1) **Temporal Mimicry**: Queries exactly 29 of 33 APIs to pass basic coverage gates, then immediately executes the backdoor target (Market Dump). Uses Projected Gradient Descent (PGD, $\varepsilon = 1.5$) to evade weight filters.
- 2) **Topology Inversion**: Queries all 33 APIs but in reversed order. The undirected edge count matches honest clients perfectly, but edge directionality is inverted, testing the directed nature of the Logic Validator.
- 3) **Gradient Mimicry (Explanation Poisoning proxy)**: Adds a FedProx proximal penalty ($\mu = 0.01$) to explicitly force its weight gradients to match the median honest update, solving the Explanation Poisoning dual objective (Equation 1). It fires the backdoor only when the VIX macro trigger spikes.

B. Ablation Study: Defense Component Contributions

Table VII presents an ablation study of the PoR defense components against the full adversary suite.

C. Topology Irreducibility and Privacy Budget

For the Temporal Mimicry attacker, querying $k = 32$ of 33 APIs before sabotaging yields a minimum theoretical Jaccard Edit Distance of $\varepsilon_{\min} = 0.0882$ —strictly above the Finance domain’s dynamic threshold of $\tau = 0.07$, mathematically guaranteeing detection regardless of weight-space camouflage.

The Finance domain implements DP-SGD with RDP composition across $T = 5$ rounds, yielding ($\varepsilon \approx 2.0, \delta = 10^{-5}$)-DP ($\sigma = 0.3, C = 1.0$). PoR filtering improves the effective privacy-utility trade-off: adversarial gradients are rejected prior to aggregation, preserving higher main task utility for the global Transformer policy while maintaining the mathematical privacy bound.

SIGNATURES

Students: _____

Project Guide: _____